

# SECOM- 1st Phase

## 1. Data Loading

Data understanding		
Purpose	Library	Query
Load dataset - SECOM	import pandas as pd	<pre># Replace the file path with the actual path to your file  file_path = "C:/Users/DhruviJayPatel/Documents/SECOMProject/seco m_data/secom.data"  # Read the data into a DataFrame  secom = pd.read_csv(file_path, header=None, sep=" ")</pre>
Load dataset Label		<pre>file_path_2 = "C:/Users/DhruviJayPatel/Documents/SECOMProject/seco m_data/secom_labels.data"  label = pd.read_csv(file_path_2, header=None, sep=" ")</pre>
Merge datasets		<pre>merged_df = pd.concat([secom, label], axis=1)  print(merged_df.columns)</pre>
Name predictors (feature1, feature2...)	import pandas as pd	<pre>import pandas as pd  # Assuming secom is your DataFrame with 591 features  # Replace secom with your actual DataFrame  # Generate new column names using a list comprehension  new_column_names = [f'feature{i}' for i in range(1, 593)]</pre>

		<pre> # Rename columns in the DataFrame  merged_df.columns = new_column_names   # Print the DataFrame to verify the changes  print(merged_df) </pre>
--	--	--

## 2. Data Understanding

Data understanding		
Purpose	Library	Query
<b>Histogram of missing values</b>	import pandas as pd import matplotlib.pyplot as plt	<pre> # Calculate the percentage of missing values for each feature missing_percentage = (secom.isnull().sum() / len(secom)) * 100   # Plot histogram of percentage of missing values plt.figure(figsize=(10, 6)) plt.hist(missing_percentage, bins=50, color='skyblue', edgecolor='black') plt.title('Histogram of Percentage of Missing Values') plt.xlabel('Percentage of Missing Values') plt.ylabel('Frequency') plt.grid(True) plt.show() </pre>

<b>Histogram of variance</b>	import pandas as pd import matplotlib.pyplot as plt import numpy as np # Add this line to import NumPy	import pandas as pd import matplotlib.pyplot as plt import numpy as np # Add this line to import NumPy # Assuming df is your DataFrame with 591 features # Replace df with your actual DataFrame  # Calculate variance for each feature variances = secom.var()  # Generate bin edges with intervals of 0.5 starting from 0  # Plot histogram of variances plt.figure(figsize=(20, 10)) plt.hist(variances, bins=50, color='skyblue', edgecolor='black') # Use bin edges generated above plt.title('Histogram of Feature Variances') plt.xlabel('Variance') plt.ylabel('Frequency') plt.grid(True) plt.show()
<b>Heatmap</b>	import pandas as pd import seaborn as sns import matplotlib.pyplot as plt	import pandas as pd import seaborn as sns import matplotlib.pyplot as plt  # Assuming df contains your DataFrame with 592 variables # Calculate the correlation matrix correlation_matrix = secom.corr()  # Plot the heatmap plt.figure(figsize=(20, 15)) sns.heatmap(correlation_matrix, cmap="coolwarm", vmin=-1, vmax=1, center=0, annot=False) plt.title('Correlation Heatmap of 592 Variables') plt.xlabel('Features') # Add x-axis label plt.ylabel('Features') # Add y-axis label plt.show()
<b>Duplicates</b>		duplicate_rows = secom.duplicated()  # Count the number of duplicate rows num_duplicates = duplicate_rows.sum()  print("Number of duplicate rows:", num_duplicates)

Duplicate_features		<pre>total_duplicate_features = sum(secom.T.duplicated())  # Print the total number of duplicate features print("Total number of duplicate features:", total_duplicate_features)</pre>
Pareto chart missing values	<pre>import pandas as pd import numpy as np import matplotlib.pyplot as plt</pre>	<pre>import pandas as pd import numpy as np import matplotlib.pyplot as plt  # Calculate the percentage of missing values for each column missing_percentage = (df.isnull().sum() / len(df)) * 100  # Sort the columns based on the percentage of missing values in descending order sorted_indices = missing_percentage.sort_values(ascending=False).index  # Create a histogram of the missing values plt.figure(figsize=(10, 6)) plt.hist(missing_percentage, bins=20, color='skyblue', edgecolor='black') plt.xlabel('Percentage of Missing Values') plt.ylabel('Frequency') plt.title('Histogram of Missing Values')  # Plot Pareto line on the same graph cumulative_percentage = (missing_percentage[sorted_indices].cumsum() / missing_percentage.sum() * 100).values plt.twinx() plt.xlim(0, 100)</pre>

## 2.1 Threshold Definition

<b>Threshold definition</b>	<pre>import matplotlib.pyplot as plt</pre>	<pre>import matplotlib.pyplot as plt  # Define the threshold for missing values (e.g., 10%) threshold = 60  # Calculate the percentage of missing values for each column in the training set train_missing_percentage = (X_train.isna().mean() * 100).round(2)  # Calculate the percentage of missing values for each column in the testing set test_missing_percentage = (X_test.isna().mean() * 100).round(2)  # Plot histogram of missing value percentages for the training set plt.figure(figsize=(12, 6)) train_hist, train_bins, _ = plt.hist(train_missing_percentage, bins=10, range=(40, 100), color='skyblue', edgecolor='black') plt.axvline(x=threshold, color='red', linestyle='--', label='Threshold')  # Add annotations for each bar in the training histogram for i, freq in enumerate(train_hist): plt.text(train_bins[i], freq, str(int(freq)), ha='center', va='bottom')  plt.title('Histogram of Missing Value Percentages in Training Set') plt.xlabel('Percentage of Missing Values') plt.ylabel('Frequency') plt.legend() plt.tight_layout() plt.show()  # Plot histogram of missing value percentages for the testing set plt.figure(figsize=(12, 6)) test_hist, test_bins, _ = plt.hist(test_missing_percentage, bins=10, range=(40, 100), color='salmon', edgecolor='black') plt.axvline(x=threshold, color='red', linestyle='--', label='Threshold')  # Add annotations for each bar in the testing histogram for i, freq in enumerate(test_hist): plt.text(test_bins[i], freq, str(int(freq)), ha='center', va='bottom')</pre>
-----------------------------	--	--

		<pre>plt.title('Histogram of Missing Value Percentages in Testing Set') plt.xlabel('Percentage of Missing Values') plt.ylabel('Frequency') plt.legend() plt.tight_layout() plt.show()</pre>
--	--	---

## 2.2 Outlier Analysis

<b>Outlier Analysis</b>	<pre>import pandas as pd import numpy as np import matplotlib.pyplot as plt</pre>	<pre>import pandas as pd import numpy as np import matplotlib.pyplot as plt  # Calculate z-scores for each feature z_scores = (X_train - X_train.mean()) / X_train.std() z_scores_test = (X_test - X_test.mean()) / X_test.std()  # Define threshold for identifying outliers threshold = 3  # Find outliers for each feature outliers = np.abs(z_scores) &gt; threshold outliers_test = np.abs(z_scores_test) &gt; threshold  # Count number of outliers for each feature num_outliers = np.sum(outliers, axis=0)  # Calculate percentage of outliers for each feature percentage_outliers = (num_outliers / len(X_train)) * 100  import matplotlib.pyplot as plt  # Define bin edges bin_edges = [0, 0.001] + list(range(1, 6))  # Plot histogram plt.figure(figsize=(10, 6)) hist = plt.hist(percentage_outliers, bins=bin_edges, color='skyblue', edgecolor='black') plt.title("Histogram of Percentages of Outliers in Each Column") plt.xlabel("Percentage of Outliers") plt.ylabel("Frequency") plt.xticks(bin_edges) for bar in hist[2]: height = int(bar.get_height()) plt.text(bar.get_x() + bar.get_width() / 2, height, height, ha='center', va='bottom') plt.grid(axis='y') plt.show()</pre>
-------------------------	---	---

### 3 Splitting

<b>Split the data</b>	<pre>from sklearn.model_selection import train_test_split</pre>	<pre>from sklearn.model_selection import train_test_split  y = merged_df['feature591'] # Replace 'target_column_name' with the name of your target column  # Dropping the target variable from the dataframe to get only the features X = merged_df.drop('feature591', axis=1)  # Split pass cases into training and testing sets while preserving the distribution X_pass_train, X_pass_test, y_pass_train, y_pass_test = train_test_split( X[y == -1], y[y == -1], test_size=0.25, random_state=42, stratify=y[y == -1])  # Split fail cases into training and testing sets while preserving the distribution X_fail_train, X_fail_test, y_fail_train, y_fail_test = train_test_split( X[y == 1], y[y == 1], test_size=0.25, random_state=42, stratify=y[y == 1])  # Concatenate the pass and fail cases in training and testing sets X_train = pd.concat([X_pass_train, X_fail_train]) y_train = pd.concat([y_pass_train, y_fail_train]) X_test = pd.concat([X_pass_test, X_fail_test]) y_test = pd.concat([y_pass_test, y_fail_test])  print("Training set - Features:", X_train.shape, "Labels:", y_train.shape) print("Testing set - Features:", X_test.shape, "Labels:", y_test.shape)</pre>
<b>Fail and Pass Proportion</b>		
<b>original dataset</b>		<pre>pass_original_proportion = (merged_df['feature591'] == -1).mean() fail_original_proportion = (merged_df['feature591'] == 1).mean()  print("Original Dataset:") print("Pass cases proportion:", pass_original_proportion) print("Fail cases proportion:", fail_original_proportion)</pre>



<b>Train and Test data</b>		<pre># Calculate the proportion of pass and fail cases in the training set pass_train_proportion = (y_train == -1).mean() fail_train_proportion = (y_train == 1).mean()  # Calculate the proportion of pass and fail cases in the testing set pass_test_proportion = (y_test == -1).mean() fail_test_proportion = (y_test == 1).mean()  print("Training Set:") print("Pass cases proportion:", pass_train_proportion) print("Fail cases proportion:", fail_train_proportion) print("\nTesting Set:") print("Pass cases proportion:", pass_test_proportion) print("Fail cases proportion:", fail_test_proportion)</pre>
----------------------------	--	--

# SECOM - 2nd Phase

## 4 Data Preparation

### 4.1 Variance

DATA PREPARATION		
Purpose	Library	Query
Removal Zero Variance Features		
Train		"variance = X_train.var() columns_to_keep = variance[variance != 0].index X_train = X_train[columns_to_keep]"
Test		X_train = X_train[columns_to_keep]"

### 4.2 Missing Value Removal

Purpose	Library	Query
Defining Missing Value Threshold and removal		"missing_percentages = X_train.isnull().mean() * 100  # Identify columns with more than 60% missing values columns_to_drop = missing_percentages[missing_percentages > 45].index  # Drop those columns from the training set X_train = X_train.drop(columns=columns_to_drop) "

## 4.3 Outlier Handling

Purpose	Library	Query
<b>Outlier Handling</b>		
<b>Remove</b>	import pandas as pd import numpy as np	<pre># Replace outliers with NaNs in X_train DataFrame X_train = X_train.mask(outliers) X_test= X_test.mask(outliers_test)  # Print the first few rows of the cleaned DataFrame to verify print("First few rows of X_train after replacing outliers with NaNs:") print(X_train.head())</pre>
<b>Replace</b>	import pandas as pd import numpy as np	<pre># Function to replace outliers based on 3 standard deviations and count them, excluding specified columns def replace_and_count_outliers(df, exclude_columns=[]):     outlier_counts = {}     for column in df.columns:         if column not in exclude_columns: mean = df[column].mean()             std = df[column].std() lower_bound =             mean - 3 * std upper_bound = mean +             3 * std              # Count outliers before replacing             outliers = ((df[column] &lt; lower_bound)   (df[column] &gt; upper_bound)).sum()             outlier_counts[column] = outliers              # Replace outliers             df[column] = np.where(df[column] &lt; lower_bound, lower_bound, df[column])             df[column] = np.where(df[column] &gt; upper_bound, upper_bound, df[column])              return df, outlier_counts  # Replace outliers in the original training set and count them X_test, outliers_count = replace_and_count_outliers(X_test)  # Convert to a DataFrame for better readability outliers_df = pd.DataFrame.from_dict(outliers_count, orient='index', columns=['Outliers Count'])  # Print the outlier counts after replacement print(outliers_df)</pre>

## 4.4 Imputation

Purpose	Library	Query
Missing Value Imputation		
Imputing	from sklearn.impute import KNNImputer	<pre>imputer = KNNImputer(n_neighbors=5)  # Fit the imputer to your data and transform it data_imputed = imputer.fit_transform(X_train)  # Convert the imputed data back to a DataFrame X_train = pd.DataFrame(data_imputed, columns=X_train.columns)  from sklearn.impute import KNNImputer  imputer = KNNImputer(n_neighbors=5)  # Fit the imputer to your data and transform it data_imputed = imputer.fit_transform(X_test)  # Convert the imputed data back to a DataFrame X_test = pd.DataFrame(data_imputed, columns=X_test.columns)</pre>

<b>Comparison</b>	<pre> import pandas as pd import matplotlib.pyplot as plt from sklearn.impute import KNNImputer from sklearn.experimental import enable_iterative_imputer from sklearn.impute import IterativeImputer, SimpleImputer </pre>	<pre> import pandas as pd import matplotlib.pyplot as plt from sklearn.impute import KNNImputer from sklearn.experimental import enable_iterative_imputer from sklearn.impute import IterativeImputer, SimpleImputer  # Calculate missing percentage per column missing_percentage = (X_train.isnull().sum() / len(data)) * 100  # Select columns with missing values between 40% to 65% cols_to_plot = missing_percentage[(missing_percentage &gt;= 40) &amp; (missing_percentage &lt;= 65)].index  # Calculate volatility metrics before imputation for selected columns volatility_before = data[cols_to_plot].std() # Using standard deviation as an example print("\nVolatility Metrics Before Imputation:") print(volatility_before)  # Imputation methods imputation_methods = ['mean', 'median', 'knn', 'mice'] imputation_results = {}  # Apply different imputation methods to selected columns only for method in imputation_methods: if method == 'mean': # Mean imputation imputer = SimpleImputer(strategy='mean') data_imputed = pd.DataFrame(imputer.fit_transform(X_train[cols_to_plot]), columns=cols_to_plot) elif method == 'median': # Median imputation imputer = SimpleImputer(strategy='median') data_imputed = pd.DataFrame(imputer.fit_transform(X_train[cols_to_plot]), columns=cols_to_plot) elif method == 'knn': # KNN imputation imputer = KNNImputer(n_neighbors=3) data_imputed = pd.DataFrame(imputer.fit_transform(X_train[cols_to_plot]), columns=cols_to_plot) elif method == 'mice': # MICE (IterativeImputer) imputation imputer = IterativeImputer() </pre>
-------------------	---	---

```

data_imputed =
pd.DataFrame(imputer.fit_transform(X_train[cols_to_plot]),
columns=cols_to_plot)

# Calculate volatility metrics after imputation for selected columns
volatility_after = data_imputed.std() # Using standard deviation as an
example
imputation_results[method] = volatility_after

# Plot before and after volatility comparison for selected columns only
plt.figure(figsize=(10, 6))
x = range(len(cols_to_plot))

plt.bar(x, volatility_before, width=0.4, alpha=0.6, color='b',
label='Before Imputation')
plt.bar(x, volatility_after, width=0.4, alpha=0.6, color='r', label='After
Imputation')

# Add labels to the bars
for i in x:
plt.text(i, volatility_before[i], f'{volatility_before[i]:.2f}', ha='center',
va='bottom', color='blue')
plt.text(i, volatility_after[i], f'{volatility_after[i]:.2f}', ha='center',
va='bottom', color='red')

plt.title(f'Volatility Comparison - {method.capitalize()} Imputation')
plt.xlabel('Features')
plt.ylabel('Standard Deviation')
plt.xticks(x, cols_to_plot, rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

# Print volatility metrics after each imputation method for selected
columns only
print("\nVolatility Metrics After Imputation:")
for method, volatility_after in imputation_results.items():
print(f"\nMethod: {method.capitalize()}")
print(volatility_after)

```

## 4.5 Feature Selection/ Reduction

Purpose	Library	Query
Scree Plot	from sklearn.preprocessing import MinMaxScaler from sklearn.decomposition import PCA from sklearn.impute import SimpleImputer import matplotlib.pyplot as plt	<pre># Scale the data using Min-Max scaling scaler = MinMaxScaler() X_train = scaler.fit_transform(X_train)  # Fit PCA pca = PCA() pca.fit(X_train)  # Plot scree plot plt.figure(figsize=(10, 6)) plt.plot(range(1, len(pca.explained_variance_ratio_) + 1), pca.explained_variance_ratio_, marker='o', linestyle='--') plt.title('Scree Plot') plt.xlabel('Number of Components') plt.ylabel('Explained Variance Ratio') plt.xticks(range(1, len(pca.explained_variance_ratio_) + 1)) plt.grid(True) plt.show()</pre>
KMO Test		<pre># Calculate the correlation matrix corr_matrix = np.corrcoef(X_train, rowvar=False) print("\nCorrelation Matrix:") print(corr_matrix)  # Calculate KMO statistic try: kmo_all, kmo_model = calculate_kmo(X_train) print(f"\nKMO statistic: {kmo_model}") except ValueError as ve: print(f"Error occurred while calculating KMO: {ve}")</pre>

<b>Relation between features and target variable</b>		<pre># Calculate correlation with target variable (assuming `target` is your target variable) correlation_with_target = X_train.corrwith(y_train)  # Sort correlation values from largest to smallest correlation_sorted = correlation_with_target.sort_values(ascending=False)  # Print the sorted correlation Series print(correlation_sorted)</pre>
<b>BORUTA</b>	<pre>import pandas as pd import sys</pre>	<pre>sys.path.append("C:/Users/DhruviJayPatel/AppData/Local/Programs/Python/Python312/Lib/site-packages") from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier  # Assuming X_train and y_train are already defined # X_train: DataFrame with your features # y_train: Series or array with your target variable  # Initialize RandomForestClassifier rf = RandomForestClassifier(n_estimators=100, random_state=42)  # Initialize Boruta boruta_selector = BorutaPy(rf, n_estimators='auto', random_state=42)  # Fit the Boruta model boruta_selector.fit(X_train.values, y_train)  # Get the boolean mask of selected features selected_features = boruta_selector.support_  # Get the column names of selected features selected_features_columns = X_train.columns[selected_features]  # Print the selected features print("Selected Features:") print(selected_features_columns)</pre>



<b>BORUTA feature ranking</b>	<pre> from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier import matplotlib.pyplot as plt import numpy as np </pre>	<pre> # Example: Initialize Boruta and fit it forest = RandomForestClassifier(n_estimators=100, random_state=42) boruta_selector = BorutaPy(forest, n_estimators='auto', random_state=42) boruta_selector.fit(X_train.values, y_train.values.ravel())  # Get feature names and Boruta rankings feature_names = data.columns boruta_rankings = boruta_selector.ranking_  # Create a list of tuples (feature, ranking) and sort by ranking features_with_ranking = list(zip(feature_names, boruta_rankings)) features_with_ranking_sorted = sorted(features_with_ranking, key=lambda x: x[1])  # Extract sorted feature names and rankings sorted_features = [feat for feat, rank in features_with_ranking_sorted] sorted_rankings = [rank for feat, rank in features_with_ranking_sorted]  # Print sorted feature rankings print("Sorted Boruta feature rankings:") for feat, rank in features_with_ranking_sorted: print(f"{feat}: {rank}") </pre>
<b>PCA</b>	<pre> from sklearn.decomposition import PCA </pre>	<pre> pca = PCA(n_components=10) # Select the number of components X_train_pca = pca.fit_transform(X_train) X_test_pca = pca.transform(X_test) </pre>

## 4.6 Balancing

Purpose	Library	Query
<b>ADASYN</b>		<pre> # Step 2: ADASYN - Handle class imbalance adasyn = ADASYN(random_state=42) X_train_resampled, y_train_resampled = adasyn.fit_resample(X_train_pca, y_train) </pre>
<b>SMOTE</b>	<pre> from imblearn.over_sampling import SMOTE </pre>	<pre> from imblearn.over_sampling import SMOTE  smote = SMOTE(random_state=42) X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train) </pre>

<b>Sampling - scatter plot</b>	<pre> import pandas as pd import numpy as np import matplotlib.pyplot as plt import seaborn as sns from sklearn.model_selection import train_test_split from sklearn.preprocessing import StandardScaler from sklearn.impute import KNNImputer from sklearn.neighbors import KNeighborsClassifier from sklearn.metrics import accuracy_score, confusion_matrix from imblearn.over_sampling import SMOTE, ADASYN, RandomOverSampler from imblearn.under_samplin g import RandomUnderSampler from collections import Counter </pre>	<pre> import pandas as pd import numpy as np import matplotlib.pyplot as plt import seaborn as sns from sklearn.model_selection import train_test_split from sklearn.preprocessing import StandardScaler from sklearn.impute import KNNImputer from sklearn.neighbors import KNeighborsClassifier from sklearn.metrics import accuracy_score, confusion_matrix from imblearn.over_sampling import SMOTE, ADASYN, RandomOverSampler from imblearn.under_sampling import RandomUnderSampler from collections import Counter  # Assuming X_train, y_train, X_test, y_test are already defined and preprocessed  # 1. Remove features with more than 65% missing values from the training data threshold = 0.65 missing_ratio = X_train.isnull().mean() features_to_drop = missing_ratio[missing_ratio &gt; threshold].index X_train.drop(features_to_drop, axis=1, inplace=True) X_test.drop(features_to_drop, axis=1, inplace=True)  # 2. Replace outliers using a 3-sigma boundary def replace_outliers(df): for col in df.select_dtypes(include=[np.number]).columns: mean = df[col].mean() std = df[col].std() upper_bound = mean + 3 * std lower_bound = mean - 3 * std df[col] = np.clip(df[col], lower_bound, upper_bound) return df  X_train = replace_outliers(X_train) X_test = replace_outliers(X_test)  # 3. Perform kNN imputation imputer = KNNImputer() X_train_imputed = imputer.fit_transform(X_train) X_test_imputed = imputer.transform(X_test)  # Convert back to DataFrame X_train = pd.DataFrame(X_train_imputed, columns=X_train.columns) </pre>
--------------------------------	---	--

```

X_test = pd.DataFrame(X_test_imputed, columns=X_test.columns)

# Function to create scatter plot
def create_scatter_plot(X, y, title, ax):
    X_sample = X.iloc[:, :2] # selecting first two features for visualization
    sns.scatterplot(x=X_sample.iloc[:, 0], y=X_sample.iloc[:, 1],
                    hue=y.map({1: 'blue', -1: 'red'}), ax=ax, palette=['blue', 'red'])
    ax.set_title(title)

# Count the number of pass and fail cases
counts = Counter(y)
pass_count = counts.get(-1, 0)
fail_count = counts.get(1, 0)

# Add text annotations to the plot
textstr = f'Majority (Pass): {pass_count}\nMinority (Fail): {fail_count}'
props = dict(boxstyle='round', facecolor='wheat', alpha=0.5)
ax.text(0.05, 0.95, textstr, transform=ax.transAxes, fontsize=12,
        verticalalignment='top', bbox=props)

# Remove legend
ax.get_legend().remove()

# Sampling methods
sampling_methods = {
    'Original': (X_train, y_train),
    'Undersampling':
        RandomUnderSampler(random_state=42).fit_resample(X_train,
        y_train),
    'Oversampling':
        RandomOverSampler(random_state=42).fit_resample(X_train,
        y_train),
    'SMOTE': SMOTE(random_state=42).fit_resample(X_train, y_train),
    'ADASYN': ADASYN(random_state=42).fit_resample(X_train,
        y_train),
    'ROSE':
        RandomOverSampler(random_state=42).fit_resample(X_train,
        y_train)
}

# Plotting scatter plots
fig, axes = plt.subplots(3, 2, figsize=(18, 18))
axes = axes.ravel()
for ax, (method, (X_resampled, y_resampled)) in zip(axes,
sampling_methods.items()):
    create_scatter_plot(X_resampled, y_resampled, method, ax)

```

```
plt.tight_layout()
plt.show()

# Training and evaluation
knn = KNeighborsClassifier(n_neighbors=5)

for method, (X_resampled, y_resampled) in
sampling_methods.items():
# Reversing class labels for training
y_resampled_reversed = np.where(y_resampled == 1, -1, 1)
knn.fit(X_resampled, y_resampled_reversed)
y_pred = knn.predict(X_test)

acc = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

print(f"Method: {method}")
print(f"Accuracy: {acc}")
print("Confusion Matrix:")
print(conf_matrix)
print("\n")

# Print the number of pass and fail cases
counts_resampled = Counter(y_resampled)
pass_count_resampled = counts_resampled.get(-1, 0)
fail_count_resampled = counts_resampled.get(1, 0)
print(f"Resampled Pass (Majority): {pass_count_resampled},
Resampled Fail (Minority): {fail_count_resampled}")
print("\n")
```

## 5. Building Basic Model

### 5.1 Model 1

```
import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.impute import KNNImputer
from boruta import BorutaPy
from imblearn.over_sampling import SMOTE

# Assuming 'X_train', 'y_train', 'X_test', 'y_test' are your training and test sets

# Step 1: Load and preprocess your data
# Example: Remove features with >65% missing values
threshold = 0.65
missing_counts = X_train.isnull().sum()
cols_to_remove = missing_counts[missing_counts / len(X_train) > threshold].index
X_train = X_train.drop(cols_to_remove, axis=1)
X_test = X_test.drop(cols_to_remove, axis=1)

# Step 2: Handle outliers using Z-score and replace with NaN
def handle_outliers_zscore(df, cols=None, threshold=3):
    if cols is None:
        cols = df.select_dtypes(include=[np.number]).columns

    for col in cols:
        z_scores = np.abs((df[col] - df[col].mean()) / df[col].std())
        df[col] = np.where(z_scores > threshold, np.nan, df[col])

    return df

# Apply outlier handling using Z-score to both X_train and X_test
X_train = handle_outliers_zscore(X_train)
X_test = handle_outliers_zscore(X_test)

# Step 3: KNN imputation to replace outliers and other missing values
knn_imputer = KNNImputer(n_neighbors=5)
X_train_imputed = pd.DataFrame(knn_imputer.fit_transform(X_train),
                                columns=X_train.columns)
X_test_imputed = pd.DataFrame(knn_imputer.transform(X_test),
                                columns=X_test.columns)

# Step 4: Feature Selection using Boruta
rf = RandomForestClassifier(n_estimators=100)
boruta_selector = BorutaPy(rf, n_estimators='auto', verbose=2, random_state=1)
boruta_selector.fit(X_train_imputed.values, y_train.values)

selected_features = X_train_imputed.columns[boruta_selector.support_]

X_train_selected = X_train_imputed[selected_features]
X_test_selected = X_test_imputed[selected_features]

# Step 5: Handling imbalanced dataset using SMOTE
smote = SMOTE(random_state=1)
```

```

X_train_balanced, y_train_balanced = smote.fit_resample(X_train_selected, y_train)

# Step 6: Train the Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=1)
rf_model.fit(X_train_balanced, y_train_balanced)

# Step 7: Model evaluation
# Predict on training set
y_train_pred = rf_model.predict(X_train_balanced)
train_accuracy = accuracy_score(y_train_balanced, y_train_pred)
train_error = 1 - train_accuracy
train_confusion_matrix = confusion_matrix(y_train_balanced, y_train_pred)

# Predict on test set
y_test_pred = rf_model.predict(X_test_selected)
test_accuracy = accuracy_score(y_test, y_test_pred)
test_error = 1 - test_accuracy
test_confusion_matrix = confusion_matrix(y_test, y_test_pred)

print("Train Confusion Matrix:")
print(train_confusion_matrix)
print("Train Accuracy:", train_accuracy)
print("Train Error:", train_error)

print("\nTest Confusion Matrix:")
print(test_confusion_matrix)
print("Test Accuracy:", test_accuracy)
print("Test Error:", test_error)

# Print model accuracy
print("\nModel Accuracy on Test Set:", test_accuracy)

```

## 5.2 Customized

```

import pandas as pd
import numpy as np
from sklearn.impute import KNNImputer
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming X_train, y_train, X_test, y_test is already defined

# 1. Remove features with more than 65% missing values
threshold = 0.65
missing_ratio = X_train.isnull().mean()
features_to_drop = missing_ratio[missing_ratio > threshold].index
X_train.drop(features_to_drop, axis=1, inplace=True)
X_test.drop(features_to_drop, axis=1, inplace=True)

# 2. Replace outliers using a 3-sigma boundary
def replace_outliers(df):

```

```
for col in df.select_dtypes(include=[np.number]).columns:
    mean = df[col].mean()
    std = df[col].std()
    upper_bound = mean + 3 * std
    lower_bound = mean - 3 * std
    df[col] = np.clip(df[col], lower_bound, upper_bound)
return df
```

```
X_train = replace_outliers(X_train)
X_test = replace_outliers(X_test)
```

```
# 3. Perform kNN imputation
imputer = KNNImputer()
X_train_imputed = imputer.fit_transform(X_train)
X_test_imputed = imputer.transform(X_test)
```

```
# Convert back to DataFrame
X_train = pd.DataFrame(X_train_imputed, columns=X_train.columns)
X_test = pd.DataFrame(X_test_imputed, columns=X_test.columns)
```

```
# 4. Select specified features for model building
selected_features = ['feature60', 'feature65', 'feature66', 'feature342', 'feature351', 'feature478',
                    'feature540', 'feature563']
X_train = X_train[selected_features]
X_test = X_test[selected_features]
```

```
# Apply SMOTE for balancing the training data
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)
```

```
# Train Random Forest model
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train_smote, y_train_smote)
```

```
# Make predictions on the test set
y_pred = rf.predict(X_test)
```

```
# Evaluate the model
acc = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
```

```
print(f"Accuracy: {acc}")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)
```

```
# Calculate and print the loss cost
def calculate_loss_cost(conf_matrix, cost_fp, cost_fn):
# Confusion matrix format:
```

```

# [[TN, FP],
# [FN, TP]]
tn, fp, fn, tp = conf_matrix.ravel()
loss_cost = (fp * cost_fp) + (fn * cost_fn)
return loss_cost

cost_fp = 1000 # Example cost for False Positive
cost_fn = 5000 # Example cost for False Negative
loss_cost = calculate_loss_cost(conf_matrix, cost_fp, cost_fn)

print(f"Loss Cost: {loss_cost}")

# Plot feature importance
feature_importances = rf.feature_importances_
indices = np.argsort(feature_importances)[::-1]

plt.figure(figsize=(12, 6))
sns.barplot(x=feature_importances[indices], y=np.array(selected_features)[indices])
plt.title('Feature Importance')
plt.xlabel('Importance')
plt.ylabel('Features')
plt.show()

```

### 5.3 Pipeline

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score,
roc_auc_score
from sklearn.decomposition import PCA
from boruta import BorutaPy
from imblearn.over_sampling import SMOTE, RandomOverSampler
from sklearn.impute import KNNImputer, SimpleImputer, IterativeImputer

# Step 1: Remove features with missing values above the threshold
thresholds = [45, 50, 55, 60, 65]

def remove_features_with_missing_values(X, threshold):
    return X.loc[:, X.isnull().mean() * 100 < threshold]

# Step 3: Handle outliers - replace with 3 standard deviation boundaries and put NA for each outlier
value
def handle_outliers(df):
    for col in df.select_dtypes(include=[np.number]).columns:
        upper_bound = df[col].mean() + 3 * df[col].std()
        lower_bound = df[col].mean() - 3 * df[col].std()
        df[col] = np.where((df[col] > upper_bound) | (df[col] < lower_bound), np.nan, df[col])
    return df

```



**# Step 4: Missing value imputation**

```
def impute_missing_values(df, method):
```

```
if method == 'KNN':
```

```
    imputer = KNNImputer()
```

```
elif method == 'MICE':
```

```
    imputer = IterativeImputer()
```

```
elif method == 'Mean':
```

```
    imputer = SimpleImputer(strategy='mean')
```

```
return pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
```

**# Step 5: Feature selection/reduction**

**# Boruta**

```
def boruta_feature_selection(X, y):
```

```
    rf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5)
```

```
    feat_selector = BorutaPy(rf, n_estimators='auto', verbose=0, random_state=42)
```

```
    feat_selector.fit(X.values, y.values)
```

```
    selected_features = X.columns[feat_selector.support_].tolist()
```

```
    return selected_features, 'Boruta'
```

**# PCA**

```
def pca_reduction(X, n_components):
```

```
    pca = PCA(n_components=n_components)
```

```
    return pca.fit_transform(X), 'PCA'
```

**# Step 6: Data sampling for imbalanced dataset**

```
def handle_imbalance(X, y, method):
```

```
if method == 'SMOTE':
```

```
    smote = SMOTE(random_state=42)
```

```
    X_res, y_res = smote.fit_resample(X, y)
```

```
elif method == 'RandomOverSampler':
```

```
    ros = RandomOverSampler(random_state=42)
```

```
    X_res, y_res = ros.fit_resample(X, y)
```

```
    return X_res, y_res
```

**# Function to calculate Loss Cost based on Confusion Matrix**

```
def calculate_loss_cost(conf_matrix, cost_fp, cost_fn):
```

```
    total_fp = conf_matrix[0, 1]
```

```
    total_fn = conf_matrix[1, 0]
```

```
    loss_cost = total_fp * cost_fp + total_fn * cost_fn
```

```
    return loss_cost
```

**# List to store model performance**

```
results = []
```

```
boruta_selected_features = [] # List to collect Boruta selected features
```

**# Iterate through thresholds**

for threshold in thresholds:

```
    X_train_thresh = remove_features_with_missing_values(X_train, threshold)
```

```

X_test_thresh = X_test[X_train_thresh.columns]

# Handle outliers
X_train_outliers = handle_outliers(X_train_thresh)
X_test_outliers = handle_outliers(X_test_thresh)

# Iterate through feature selection/reduction methods

for feature_method_choice in ['Boruta', 'PCA']:
    # Iterate through imputation methods
    for impute_method in ['KNN', 'MICE', 'Mean']:
        X_train_imputed = impute_missing_values(X_train_outliers, impute_method)
        X_test_imputed = impute_missing_values(X_test_outliers, impute_method)

        # Perform feature selection/reduction based on choice
        if feature_method_choice == 'Boruta':
            selected_features, feature_method = boruta_feature_selection(X_train_imputed, y_train)
            X_train_selected = X_train_imputed[selected_features]
            X_test_selected = X_test_imputed[selected_features]
        elif feature_method_choice == 'PCA':
            X_train_selected, feature_method = pca_reduction(X_train_imputed, n_components=30)
            X_test_selected = pca_reduction(X_test_imputed, n_components=30)[0]

        # Iterate through imbalance methods
        for imbalance_method in ['SMOTE', 'RandomOverSampler']:
            X_train_res, y_train_res = handle_imbalance(X_train_selected, y_train, imbalance_method)

            # Train Random Forest model
            rf_model = RandomForestClassifier(random_state=42)
            rf_model.fit(X_train_res, y_train_res)

            # Predictions
            y_pred = rf_model.predict(X_test_selected)

            # Evaluation metrics
            accuracy = accuracy_score(y_test, y_pred)
            conf_matrix = confusion_matrix(y_test, y_pred)
            precision = precision_score(y_test, y_pred)
            recall = recall_score(y_test, y_pred)
            f1 = f1_score(y_test, y_pred)
            auc = roc_auc_score(y_test, y_pred)

            # Calculate Loss Cost
            cost_fp = 1000 # Example cost for False Positive
            cost_fn = 5000 # Example cost for False Negative
            loss_cost = calculate_loss_cost(conf_matrix, cost_fp, cost_fn)

            # Prepare results
            result = {
                'threshold': threshold,
                'impute_method': impute_method,

```

```

'imbalance_method': imbalance_method,
'accuracy': accuracy,
'confusion_matrix': conf_matrix,
'precision': precision,
'recall': recall,
'f1_score': f1,
'auc': auc,
'loss_cost': loss_cost,
'features_used': feature_method,
'selected_features': selected_features if feature_method == 'Boruta' else None
}

```

```

results.append(result)

```

```

# Collect Boruta selected features
if feature_method == 'Boruta':
    boruta_selected_features.append(selected_features)

```

```

# Convert results to DataFrame
results_df = pd.DataFrame(results)

```

```

# Display top 10 models by loss cost
top_10_models_loss_cost = results_df.sort_values(by='loss_cost', ascending=True).head(10)
print("\nTop 10 Models by Loss Cost:")
print(top_10_models_loss_cost[['threshold', 'impute_method', 'imbalance_method', 'accuracy',
'confusion_matrix',
'precision', 'recall', 'f1_score', 'auc', 'loss_cost', 'features_used', 'selected_features']])

```

```

# Display top 10 models by minimizing False Positive and False Negative errors
top_10_models_fp_fn = results_df.sort_values(by=['confusion_matrix'], key=lambda x: x.apply(lambda
y: (y[0][1], y[1][0])), ascending=True).head(10)
print("\nTop 10 Models by False Positive and False Negative Errors:")
print(top_10_models_fp_fn[['threshold', 'impute_method', 'imbalance_method', 'accuracy',
'confusion_matrix',
'precision', 'recall', 'f1_score', 'auc', 'loss_cost', 'features_used', 'selected_features']])

```

```

# Print Boruta selected features if present in top models by loss cost
if boruta_selected_features:
    print("\nBoruta Selected Features in Top Models:")
    for idx, features in enumerate(boruta_selected_features, start=1):
        print(f"Iteration {idx}: {features}")
    else:
        print("\nNo Boruta selected features in Top Models."

```

# SECOM – 3rd Phase

## 6 Modeling and Evaluation

### 6.1 Scaling

Purpose	Library	Query
Scaling	import numpy as np import matplotlib.pyplot as plt from sklearn.preprocessing import MinMaxScaler from scipy import stats	<pre># Assuming 'merged_df' is your DataFrame with 'feature17' feature17 = merged_df['feature17'].values.reshape(-1, 1)  # Apply different transformations  # Min-Max Scaling scaler_minmax = MinMaxScaler() feature17_minmax = scaler_minmax.fit_transform(feature17)  # Log Transformation feature17_log = np.log(feature17 + 1) # Adding 1 to handle zeros  # Linear Transformation feature17_linear = 0.5 * feature17 # Scaling by a factor of 0.5  # Box-Cox Transformation feature17_boxcox, _ = stats.boxcox(feature17.flatten() + 1) # Adding 1 to handle zeros  # Plotting fig, axes = plt.subplots(2, 2, figsize=(18, 12))  # Min-Max Scaling axes[0, 0].hist(feature17_minmax.flatten(), bins=30, color='green', alpha=0.7) axes[0, 0].set_title('Min-Max Scaling (Feature 17)')  # Log Transformation axes[0, 1].hist(feature17_log.flatten(), bins=30, color='purple', alpha=0.7) axes[0, 1].set_title('Log Transformation (Feature 17)')  # Linear Transformation axes[1, 0].hist(feature17_linear.flatten(), bins=30, color='orange', alpha=0.7) axes[1, 0].set_title('Linear Transformation (Feature 17)')  # Box-Cox Transformation axes[1, 1].hist(feature17_boxcox, bins=30, color='red', alpha=0.7) axes[1, 1].set_title('Box-Cox Transformation (Feature 17)')  plt.tight_layout() plt.show()</pre>

## 6.2 Optimal Parameters

Purpose	Library	Query
<b>Optimal Parameters</b>	import numpy as np import pandas as pd import matplotlib.pyplot as plt from sklearn.decomposition import PCA from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE, ADASYN from imblearn.combine import SMOTETomek from sklearn.tree import DecisionTreeClassifier from sklearn.linear_model import LogisticRegression from sklearn.naive_bayes import GaussianNB from sklearn.svm import SVC from sklearn.neighbors import KNeighborsClassifier from sklearn.preprocessing import StandardScaler from sklearn.metrics import f1_score, precision_score, recall_score, confusion_matrix	<pre># Define models models = { 'RandomForest': RandomForestClassifier(), 'DecisionTree': DecisionTreeClassifier(), 'LogisticRegression': LogisticRegression(), 'NaiveBayes': GaussianNB(), 'SVM': SVC(), 'KNN': KNeighborsClassifier() }  # Define feature selection and balancing techniques feature_methods = ['PCA', 'Boruta'] balance_methods = ['SMOTE', 'ADASYN', 'ROSE']  # Function to apply PCA def apply_pca(X_train, X_test, n_components=10): pca = PCA(n_components=n_components) X_train_pca = pca.fit_transform(X_train) X_test_pca = pca.transform(X_test) return X_train_pca, X_test_pca  # Function to apply Boruta def apply_boruta(X_train, y_train, X_test): rf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5) boruta = BorutaPy(rf, n_estimators='auto', verbose=0, random_state=1) boruta.fit(X_train.values, y_train) X_train_boruta = boruta.transform(X_train.values) X_test_boruta = boruta.transform(X_test.values) return X_train_boruta, X_test_boruta  # Balancing functions def balance_data(X_train, y_train, method): if method == 'SMOTE': smote = SMOTE(random_state=42) X_res, y_res = smote.fit_resample(X_train, y_train) elif method == 'ADASYN': adasyn = ADASYN(random_state=42) X_res, y_res = adasyn.fit_resample(X_train, y_train) elif method == 'ROSE': rose = SMOTETomek(random_state=42) X_res, y_res = rose.fit_resample(X_train, y_train) return X_res, y_res  # Evaluation function def evaluate_model(model, X_train, y_train, X_test, y_test): model.fit(X_train, y_train) y_pred = model.predict(X_test) f1 = f1_score(y_test, y_pred)</pre>

```

precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
loss = fp * 1000 + fn * 5000
return f1, precision, recall, loss

# Collecting results
results = []
for feature_method in feature_methods:
    if feature_method == 'PCA':
        X_train_fs, X_test_fs = apply_pca(X_train_imputed,
        X_test_imputed)
    elif feature_method == 'Boruta':
        X_train_fs, X_test_fs = apply_boruta(X_train_imputed, y_train,
        X_test_imputed)

for balance_method in balance_methods:
    X_train_bal, y_train_bal = balance_data(X_train_fs, y_train,
    balance_method)

# Apply scaling where necessary
scaler = StandardScaler()
X_train_bal_scaled = scaler.fit_transform(X_train_bal)
X_test_fs_scaled = scaler.transform(X_test_fs)

for model_name, model in models.items():
    if model_name in ['SVM', 'KNN', 'LogisticRegression']:
        f1, precision, recall, loss = evaluate_model(model,
        X_train_bal_scaled, y_train_bal, X_test_fs_scaled, y_test)
    else:
        f1, precision, recall, loss = evaluate_model(model, X_train_bal,
        y_train_bal, X_test_fs, y_test)
    results.append([f'{feature_method}+{balance_method}',
    model_name, f1, precision, recall, loss])

# Convert results to DataFrame
results_df = pd.DataFrame(results, columns=['Method', 'Model',
'F1 Score', 'Precision', 'Recall', 'Loss'])

# Plotting
fig, axs = plt.subplots(2, 2, figsize=(15, 10))

# Plot F1 Score
x = np.arange(len(results_df['Method'].unique()))
width = 0.15

for i, model_name in enumerate(models.keys()):
    subset = results_df[results_df['Model'] == model_name]
    axs[0, 0].bar(x + i*width, subset['F1 Score'], width,
    label=model_name)
    axs[0, 0].set_title('F1 Score')
    axs[0, 0].set_xticks(x + width*(len(models)/2))
    axs[0, 0].set_xticklabels(results_df['Method'].unique(),
    rotation=90)
    axs[0, 0].legend()

# Plot Precision
for i, model_name in enumerate(models.keys()):

```

```

subset = results_df[results_df['Model'] == model_name]
axs[0, 1].bar(x + i*width, subset['Precision'], width,
label=model_name)
axs[0, 1].set_title('Precision')
axs[0, 1].set_xticks(x + width*(len(models)/2))
axs[0, 1].set_xticklabels(results_df['Method'].unique(),
rotation=90)
axs[0, 1].legend()

# Plot Recall
for i, model_name in enumerate(models.keys()):
subset = results_df[results_df['Model'] == model_name]
axs[1, 0].bar(x + i*width, subset['Recall'], width,
label=model_name)
axs[1, 0].set_title('Recall')
axs[1, 0].set_xticks(x + width*(len(models)/2))
axs[1, 0].set_xticklabels(results_df['Method'].unique(),
rotation=90)
axs[1, 0].legend()

# Plot Loss
for i, model_name in enumerate(models.keys()):
subset = results_df[results_df['Model'] == model_name]
axs[1, 1].bar(x + i*width, subset['Loss'], width,
label=model_name)
axs[1, 1].set_title('Loss')
axs[1, 1].set_xticks(x + width*(len(models)/2))
axs[1, 1].set_xticklabels(results_df['Method'].unique(),
rotation=90)
axs[1, 1].legend()

plt.tight_layout()
plt.show()

```

## 6.3 Finding best fit models among RF, NB and SVM

Modeling and Evaluation		
Purpose	Library	Query
<b>Hyperparameter tuning + graphs</b>	import pandas as pd import numpy as np from sklearn.model_selection import train_test_split, GridSearchCV, KFold from sklearn.impute import KNNImputer from sklearn.preprocessing import StandardScaler from sklearn.ensemble import RandomForestClassifier from sklearn.svm import SVC from sklearn.naive_bayes import GaussianNB from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_curve, auc from imblearn.over_sampling import SMOTE import matplotlib.pyplot as plt from boruta import BorutaPy	# Selecting specific features identified by Boruta selected_features = ['feature60', 'feature65', 'feature66', 'feature342', 'feature351', 'feature478', 'feature540', 'feature563']  X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features]  # SMOTE Data Balancing sm = SMOTE(random_state=42) X_train_res, y_train_res = sm.fit_resample(X_train_selected, y_train)  # Scaling the Features for SVM only scaler = StandardScaler() X_train_res_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected)  # Function to evaluate model and return evaluation metrics def evaluate_model(model, X_train, y_train, X_test, y_test): model.fit(X_train, y_train) y_pred = model.predict(X_test) y_pred_prob = model.predict_proba(X_test)[:, 1]  accuracy = accuracy_score(y_test, y_pred) precision = precision_score(y_test, y_pred) recall = recall_score(y_test, y_pred) f1 = f1_score(y_test, y_pred) conf_matrix = confusion_matrix(y_test, y_pred)  return accuracy, precision, recall, f1, conf_matrix, y_pred_prob  # Define models models = [ SVC(probability=True), GaussianNB(), RandomForestClassifier(n_estimators=100, random_state=42) ]  # Define colors for each model consistently model_colors = { 'RandomForestClassifier': 'green', 'GaussianNB': 'lightgreen', 'SVC': 'red' }  # Lists to store results before and after tuning models_results_before = [] models_results_after = [] roc_curves_before = [] roc_curves_after = []  # Evaluate each model before tuning for model in models:



```

model_name = model.__class__.__name__
color = model_colors[model_name]

if isinstance(model, SVC):
    accuracy, precision, recall, f1, conf_matrix, y_pred_prob =
    evaluate_model(model, X_train_res_scaled, y_train_res,
    X_test_scaled, y_test)
else:
    accuracy, precision, recall, f1, conf_matrix, y_pred_prob =
    evaluate_model(model, X_train_res, y_train_res, X_test_selected,
    y_test)

models_results_before.append({
'Model': model_name,
'Accuracy': accuracy,
'Precision': precision,
'Recall': recall,
'F1 Score': f1,
'Confusion Matrix': conf_matrix,
'Color': color
})

fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)
roc_curves_before.append((fpr, tpr, roc_auc, model_name, color))

# Define parameter grids for tuning
param_grids = {
'SVC': {'C': [0.1, 1, 10], 'gamma': ['scale', 'auto'], 'kernel': ['linear',
'rbf']},
'GaussianNB': {'var_smoothing': [1e-09, 1e-08, 1e-07]},
'RandomForestClassifier': {'max_depth': [None, 5, 10]}
}

# Define k-fold Cross-validation
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Evaluate each model after tuning with k-fold
for model in models:
    model_name = model.__class__.__name__
    param_grid = param_grids[model_name]
    color = model_colors[model_name]

    grid_search = GridSearchCV(model, param_grid, scoring='roc_auc',
    cv=kf, n_jobs=-1)
    if isinstance(model, SVC):
        grid_search.fit(X_train_res_scaled, y_train_res)
        best_model = grid_search.best_estimator_
        accuracy, precision, recall, f1, conf_matrix, y_pred_prob =
        evaluate_model(best_model, X_train_res_scaled, y_train_res,
        X_test_scaled, y_test)
    else:
        grid_search.fit(X_train_res, y_train_res)
        best_model = grid_search.best_estimator_
        accuracy, precision, recall, f1, conf_matrix, y_pred_prob =
        evaluate_model(best_model, X_train_res, y_train_res,
        X_test_selected, y_test)

```

```

best_params = grid_search.best_params_
best_score = grid_search.best_score_

print(f"Best Parameters for {model_name}: {best_params}")
print(f"Best ROC AUC Score: {best_score}")

models_results_after.append({
    'Model': model_name,
    'Accuracy': accuracy,
    'Precision': precision,
    'Recall': recall,
    'F1 Score': f1,
    'Confusion Matrix': conf_matrix,
    'Color': color
})

fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)
roc_curves_after.append((fpr, tpr, roc_auc, model_name, color))

# Plotting comparison graphs
metrics = ['Precision', 'Recall', 'F1 Score', 'Accuracy']
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

for i, metric in enumerate(metrics):
    ax = axes[i//2, i%2]
    index = np.arange(len(models_results_before))
    bar_width = 0.35

    before_values = [result[metric] for result in models_results_before]
    after_values = [result[metric] for result in models_results_after]
    colors = [result['Color'] for result in models_results_before]

    bars1 = ax.bar(index, before_values, bar_width, color=colors,
        alpha=0.6)
    bars2 = ax.bar(index + bar_width, after_values, bar_width,
        color=colors, alpha=1.0)

    ax.set_xlabel('Models')
    ax.set_ylabel(metric)
    ax.set_title(f'{metric} Comparison')
    ax.set_xticks(index + bar_width / 2)
    ax.set_xticklabels([result['Model'] for result in
        models_results_before])

    for bar in bars1:
        yval = bar.get_height()
        ax.text(bar.get_x() + bar.get_width() / 2, yval / 2, f'{yval:.2f}',
            ha='center', va='center', color='white')

    for bar in bars2:
        yval = bar.get_height()
        ax.text(bar.get_x() + bar.get_width() / 2, yval / 2, f'{yval:.2f}',
            ha='center', va='center', color='white')

fig.tight_layout()

# Add single legend for all subplots

```

```

handles = [plt.Rectangle((0,0),1,1, color='gray', alpha=0.6,
label='Before Tuning'),
plt.Rectangle((0,0),1,1, color='gray', alpha=1.0, label='After Tuning')]
fig.legend(handles=handles, loc='upper center',
bbox_to_anchor=(0.5, 1.05), ncol=2)
plt.show()

# Plotting TP, TN, FP, FN side by side with labels
fig, axes = plt.subplots(1, 4, figsize=(18, 6))
metrics_names = ['TN', 'FN', 'FP', 'TP']

# Confusion matrix plotting function
def plot_conf_matrix(metric_name, i, ax):
index = np.arange(len(models_results_before))
bar_width = 0.35

before_values = [result['Confusion Matrix'].ravel()[i] for result in
models_results_before]
after_values = [result['Confusion Matrix'].ravel()[i] for result in
models_results_after]
colors = [result['Color'] for result in models_results_before]

bars1 = ax.bar(index, before_values, bar_width, color=colors,
alpha=0.6)
bars2 = ax.bar(index + bar_width, after_values, bar_width,
color=colors, alpha=1.0)

ax.set_xlabel('Models')
ax.set_ylabel('Counts')
ax.set_title(f'{metric_name} Comparison')
ax.set_xticks(index + bar_width / 2)
ax.set_xticklabels([result['Model'] for result in
models_results_before])

for bar in bars1:
yval = bar.get_height()
ax.text(bar.get_x() + bar.get_width() / 2, yval / 2, f'{yval:.0f}',
ha='center', va='center', color='white')

for bar in bars2:
yval = bar.get_height()
ax.text(bar.get_x() + bar.get_width() / 2, yval / 2, f'{yval:.0f}',
ha='center', va='center', color='white')

for i, metric_name in enumerate(metrics_names):
plot_conf_matrix(metric_name, i, axes[i])

fig.tight_layout()
plt.show()

# Plot ROC Curves
plt.figure(figsize=(10, 8))
for fpr, tpr, roc_auc, model_name, color in roc_curves_before:
plt.plot(fpr, tpr, lw=2, linestyle='--', color=color, label=f'{model_name}
Before (AUC = {roc_auc:.2f})')

for fpr, tpr, roc_auc, model_name, color in roc_curves_after:
plt.plot(fpr, tpr, lw=2, color=color, label=f'{model_name} After (AUC

```

```

= {roc_auc:.2f}'))

plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')

# Create a single legend outside the plot for ROC Curves
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()

# Function to calculate loss cost from confusion matrix
def calculate_loss_cost(conf_matrix):
# Define the costs for each type of misclassification
cost_fp = 1000 # Cost of False Positive
cost_fn = 5000 # Cost of False Negative

# Extract values from confusion matrix
FP = conf_matrix[0, 1]
FN = conf_matrix[1, 0]

# Calculate total loss cost
total_cost = FP * cost_fp + FN * cost_fn
return total_cost

# Extracting model names and corresponding colors
model_names = [result['Model'] for result in models_results_before]
colors = [result['Color'] for result in models_results_before]

# Loss costs before and after tuning
loss_costs_before = [calculate_loss_cost(result['Confusion Matrix'])
for result in models_results_before]
loss_costs_after = [calculate_loss_cost(result['Confusion Matrix']) for
result in models_results_after]

# Setting up the figure
fig, ax = plt.subplots(figsize=(10, 6))
bar_width = 0.35
index = np.arange(len(model_names))

bars1 = ax.bar(index - bar_width/2, loss_costs_before, bar_width,
color=colors, alpha=0.6, label='Before Tuning')
bars2 = ax.bar(index + bar_width/2, loss_costs_after, bar_width,
color=colors, alpha=1.0, label='After Tuning')

ax.set_xlabel('Models')
ax.set_ylabel('Loss Cost')
ax.set_title('Loss Cost Comparison Before and After Tuning')
ax.set_xticks(index)
ax.set_xticklabels(model_names)
ax.legend()

# Adding text labels for values on top of bars
def autolabel(bars):

```

		<pre> for bar in bars:     yval = bar.get_height()     ax.text(bar.get_x() + bar.get_width()/2, yval / 2, f'{yval:.0f}',             ha='center', va='center', color='white')  autolabel(bars1) autolabel(bars2)  plt.tight_layout() plt.show() </pre>
<b>Learning curve</b>	<pre> from sklearn.model_selection import cross_val_score, learning_curve import matplotlib.pyplot as plt </pre>	<pre> kf = KFold(n_splits=5, shuffle=True, random_state=42)  # Cross-validation scores def cross_val_evaluation(model, X, y, cv=kf):     accuracy = cross_val_score(model, X, y, cv=cv, scoring='accuracy')     precision = cross_val_score(model, X, y, cv=cv, scoring='precision')     recall = cross_val_score(model, X, y, cv=cv, scoring='recall')     f1 = cross_val_score(model, X, y, cv=cv, scoring='f1')     return accuracy, precision, recall, f1  # Learning curves def plot_learning_curve(estimator, title, X, y, cv=kf):     plt.figure()     plt.title(title)     plt.xlabel("Training examples")     plt.ylabel("Score")     train_sizes, train_scores, test_scores = learning_curve(estimator, X, y, cv=cv, n_jobs=-1)     train_scores_mean = np.mean(train_scores, axis=1)     test_scores_mean = np.mean(test_scores, axis=1)     plt.grid()      plt.plot(train_sizes, train_scores_mean, 'o-', color="r", label="Training score")     plt.plot(train_sizes, test_scores_mean, 'o-', color="g", label="Cross- validation score")      plt.legend(loc="best")     plt.show()  # usage for cross-validation models = [     (SVC(probability=True), X_train_res_scaled, y_train_res),     (GaussianNB(), X_train_res, y_train_res),     (RandomForestClassifier(n_estimators=100, random_state=42), X_train_res, y_train_res) ]  for model, X, y in models:     acc, prec, rec, f1 = cross_val_evaluation(model, X, y)     print(f"{model.__class__.__name__} - Accuracy: {np.mean(acc):.2f}, Precision: {np.mean(prec):.2f}, Recall: {np.mean(rec):.2f}, F1 Score: {np.mean(f1):.2f}")  # Example usage for learning curves plot_learning_curve(SVC(probability=True), "Learning Curve (SVM)", X_train_res_scaled, y_train_res) </pre>

		<pre> plot_learning_curve(GaussianNB(), "Learning Curve (Naive Bayes)", X_train_res, y_train_res) plot_learning_curve(RandomForestClassifier(n_estimators=100, random_state=42), "Learning Curve (Random Forest)", X_train_res, y_train_res) </pre>
--	--	---

## 6.4 Feature Engineering

Modeling and Evaluation		
Purpose	Library	Query
Feature Engineering	import pandas as pd import numpy as np	<pre> # Assuming 'merged_df' is your DataFrame and 'feature592' is the column with date-time values  # Convert the 'feature592' column to datetime format merged_df['feature592'] = pd.to_datetime(merged_df['feature592'], format='%d/%m/%Y %H:%M:%S')  # Extract date and time components merged_df['year'] = merged_df['feature592'].dt.year merged_df['month'] = merged_df['feature592'].dt.month merged_df['day'] = merged_df['feature592'].dt.day merged_df['hour'] = merged_df['feature592'].dt.hour merged_df['minute'] = merged_df['feature592'].dt.minute merged_df['second'] = merged_df['feature592'].dt.second  # Day of the week (0=Monday, 6=Sunday) merged_df['day_of_week'] = merged_df['feature592'].dt.dayofweek  # Part of the day (morning, afternoon, evening, night) def part_of_day(hour):     if 5 &lt;= hour &lt; 12:         return 1 # morning     elif 12 &lt;= hour &lt; 17:         return 2 # afternoon     elif 17 &lt;= hour &lt; 21:         return 3 # evening     else:         return 0 # night  merged_df['part_of_day'] = merged_df['hour'].apply(part_of_day)  # Elapsed time since first entry in minutes merged_df['elapsed_time'] = (merged_df['feature592'] - merged_df['feature592'].min()).dt.total_seconds() / 60  # Seasonal features merged_df['quarter'] = merged_df['feature592'].dt.quarter  # Flag for weekend merged_df['is_weekend'] = (merged_df['day_of_week'] &gt;= 5).astype(int)  # Time differences between consecutive entries in minutes merged_df['time_diff'] = merged_df['feature592'].diff().dt.total_seconds() / 60 merged_df['time_diff'].fillna(0, inplace=True) </pre>

		<pre># Display the DataFrame with new features print(merged_df.head())  # Drop the original datetime column if not needed merged_df.drop('feature592', axis=1, inplace=True)</pre>
--	--	--

## 6.5 Building models after Feature Engineering and Evaluation

Modeling and Evaluation		
Purpose	Library	Query
<b>Before tuning-Model 1 (RF with General Boruta)</b>	<pre>import pandas as pd import numpy as np from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score</pre>	<pre># Assuming 'merged_df' is your DataFrame with the relevant data X = merged_df.drop('feature591', axis=1) # Replace 'feature591' with your actual target column name y = merged_df['feature591']  # Split data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)  # Boruta Feature Selection rf_clf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5, random_state=42) boruta_selector = BorutaPy(rf_clf, n_estimators='auto', verbose=0, random_state=42) # Set verbose to 0 boruta_selector.fit(X_train.values, y_train)  # Selected features selected_features = X_train.columns[boruta_selector.support_].tolist() print("Selected Features: ", selected_features)  # Update X_train and X_test with selected features X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features]  # SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected)  # Train Random Forest Classifier rf_clf.fit(X_train_scaled, y_train_res) y_pred = rf_clf.predict(X_test_scaled)  # Confusion Matrix conf_matrix = confusion_matrix(y_test, y_pred) tn, fp, fn, tp = conf_matrix.ravel()  # Calculate Metrics</pre>

		<pre> accuracy = accuracy_score(y_test, y_pred) f1 = f1_score(y_test, y_pred, average='weighted') precision = precision_score(y_test, y_pred, average='weighted') recall = recall_score(y_test, y_pred, average='weighted') specificity = tn / (tn + fp)  # Calculate Loss Cost cost_fp = 1000 cost_fn = 5000 loss_cost = cost_fp * fp + cost_fn * fn  print("Accuracy Score: ", accuracy) print("F1 Score: ", f1) print("Precision: ", precision) print("Recall (Sensitivity): ", recall) print("Specificity: ", specificity) print("Loss Cost: ", loss_cost)  print("Confusion Matrix:") print(conf_matrix) print(f"True Positives (TP): {tp}") print(f"True Negatives (TN): {tn}") print(f"False Positives (FP): {fp}") print(f"False Negatives (FN): {fn}") </pre>
<b>Before tuning- Model 2 (RF selected features)</b>	<pre> import pandas as pd import numpy as np from sklearn.model_selection import train_test_split from sklearn.preprocessing import MinMaxScaler from sklearn.ensemble import RandomForestClassifier from sklearn.neighbors import KNeighborsClassifier from sklearn.naive_bayes import GaussianNB from sklearn.tree import DecisionTreeClassifier from sklearn.linear_model import LogisticRegression from sklearn.svm import SVC from imblearn.over_sampling import SMOTE from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix import matplotlib.pyplot as plt </pre>	<pre> # Assuming 'merged_df' is your DataFrame with the selected features and target selected_features = ['feature17', 'feature41', 'feature60', 'feature66', 'feature342', 'feature427', 'feature442', 'feature563', 'elapsed_time'] X = merged_df[selected_features] y = merged_df['feature591'] # Replace 'feature591' with your actual target column name  # Min-Max scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train) X_test_scaled = scaler.transform(X_test)  # SMOTE Data Balancing smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_scaled, y_train)  # Define classifiers classifiers = {     'Random Forest': RandomForestClassifier(random_state=42),     'KNN': KNeighborsClassifier(),     'Naive Bayes': GaussianNB(),     'Decision Tree': DecisionTreeClassifier(random_state=42),     'Logistic Regression': LogisticRegression(random_state=42),     'SVM': SVC(probability=True, random_state=42) }  # Dictionary to store evaluation metrics metrics = {     'Accuracy': [],     'Precision': [],     'Recall': [], </pre>



```

'F1 Score': [],
'ROC AUC': [],
'Loss Cost': [],
'Confusion Matrix': []
}

# Iterate over classifiers
for clf_name, clf in classifiers.items():
# Train the model
clf.fit(X_train_res, y_train_res)

# Predictions
y_pred = clf.predict(X_test_scaled)
y_pred_prob = clf.predict_proba(X_test_scaled)[:, 1]

# Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_prob)
conf_matrix = confusion_matrix(y_test, y_pred)

# Calculate loss cost
cost_fp = 1000
cost_fn = 5000
FP = conf_matrix[0, 1]
FN = conf_matrix[1, 0]
loss_cost = cost_fp * FP + cost_fn * FN

# Store metrics
metrics['Accuracy'].append(accuracy)
metrics['Precision'].append(precision)
metrics['Recall'].append(recall)
metrics['F1 Score'].append(f1)
metrics['ROC AUC'].append(roc_auc)
metrics['Loss Cost'].append(loss_cost)
metrics['Confusion Matrix'].append(conf_matrix)

# Print results
print(f"Classifier: {clf_name}")
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
print(f"ROC AUC: {roc_auc}")
print(f"Confusion Matrix:\n{conf_matrix}")
print(f"Loss Cost: {loss_cost}")
print("\n")

# Plotting ROC curves
plt.figure(figsize=(8, 6))
for clf_name, clf in classifiers.items():
y_pred_prob = clf.predict_proba(X_test_scaled)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, label=f'{clf_name} (AUC = {roc_auc:.2f})')

```

		<pre> plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--') plt.xlim([0.0, 1.0]) plt.ylim([0.0, 1.05]) plt.xlabel('False Positive Rate') plt.ylabel('True Positive Rate') plt.title('Receiver Operating Characteristic') plt.legend(loc="lower right") plt.show() </pre>
<b>Before tuning- Model 3 (Gaussian NB)</b>	<pre> import pandas as pd import numpy as np from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score from sklearn.naive_bayes import GaussianNB </pre>	<pre> # Assuming X_train_imputed, y_train, X_test_imputed, y_test are already defined X_train = X_train_imputed[['feature17', 'feature41', 'feature60', 'feature66', 'feature342', 'feature427', 'feature442', 'feature563', 'elapsed_time']].copy() y_train = y_train.copy() X_test = X_test_imputed[['feature17', 'feature41', 'feature60', 'feature66', 'feature342', 'feature427', 'feature442', 'feature563', 'elapsed_time']].copy() y_test = y_test.copy()  # Update X_train and X_test with selected features X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features]  # SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected)  # Define parameter grid for GaussianNB param_grid = { 'var_smoothing': np.logspace(0,-9, num=100) }  # Perform parameter tuning best_f1_score = -1 best_params = None  for var_smoothing in param_grid['var_smoothing']: # Train GaussianNB with current parameters nb_clf = GaussianNB(var_smoothing=var_smoothing) nb_clf.fit(X_train_scaled, y_train_res)  # Predict on test data y_pred = nb_clf.predict(X_test_scaled)  # Calculate F1 score f1 = f1_score(y_test, y_pred, average='weighted')  # Check if current model is better than previous ones if f1 &gt; best_f1_score: best_f1_score = f1 best_params = { 'var_smoothing': var_smoothing } </pre>

		<pre> # Train final GaussianNB with best parameters best_nb_clf = GaussianNB(var_smoothing=best_params['var_smoothing']) best_nb_clf.fit(X_train_scaled, y_train_res)  # Predict on test data with the best model y_pred = best_nb_clf.predict(X_test_scaled)  # Confusion Matrix conf_matrix = confusion_matrix(y_test, y_pred) tn, fp, fn, tp = conf_matrix.ravel()  # Calculate Metrics accuracy = accuracy_score(y_test, y_pred) f1 = f1_score(y_test, y_pred, average='weighted') precision = precision_score(y_test, y_pred, average='weighted') recall = recall_score(y_test, y_pred, average='weighted') specificity = tn / (tn + fp)  # Calculate Loss Cost cost_fp = 1000 cost_fn = 5000 loss_cost = cost_fp * fp + cost_fn * fn  print("Best Parameters:", best_params) print("Best F1 Weighted Score:", best_f1_score)  print("Accuracy Score: ", accuracy) print("F1 Score: ", f1) print("Precision: ", precision) print("Recall (Sensitivity): ", recall) print("Specificity: ", specificity) print("Loss Cost: ", loss_cost)  print("Confusion Matrix:") print(conf_matrix) print(f"True Positives (TP): {tp}") print(f"True Negatives (TN): {tn}") print(f"False Positives (FP): {fp}") print(f"False Negatives (FN): {fn}") </pre>
<b>After tuning- Model 1</b>	<pre> import pandas as pd import numpy as np from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, </pre>	<pre> # Assuming 'merged_df' is your DataFrame with the relevant data X = merged_df.drop('feature591', axis=1) # Replace 'feature591' with your actual target column name y = merged_df['feature591']  # Split data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)  # Boruta Feature Selection rf_clf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5, random_state=42) boruta_selector = BorutaPy(rf_clf, n_estimators='auto', verbose=0, random_state=42) # Set verbose to 0 boruta_selector.fit(X_train.values, y_train) </pre>

<pre>precision_score, recall_score from sklearn.naive_bayes import GaussianNB</pre>	<pre># Selected features selected_features = X_train.columns[boruta_selector.support_].tolist() print("Selected Features: ", selected_features)  # Update X_train and X_test with selected features X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features]  # SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected)  # Define parameter grid for RandomForestClassifier param_grid = { 'n_estimators': [50, 100, 200, 300], 'max_depth': [3, 5, 7, 10, 15], 'min_samples_split': [2, 5, 10, 15], 'min_samples_leaf': [1, 2, 4, 6] }  # Perform hyperparameter tuning without cross-validation best_f1_score = -1 best_params = None  for n_estimators in param_grid['n_estimators']: for max_depth in param_grid['max_depth']: for min_samples_split in param_grid['min_samples_split']: for min_samples_leaf in param_grid['min_samples_leaf']: # Train RandomForestClassifier with current parameters rf_clf = RandomForestClassifier(n_estimators=n_estimators, max_depth=max_depth, min_samples_split=min_samples_split, min_samples_leaf=min_samples_leaf, n_jobs=-1, class_weight='balanced', random_state=42)  rf_clf.fit(X_train_scaled, y_train_res)  # Predict on test data y_pred = rf_clf.predict(X_test_scaled)  # Calculate F1 score f1 = f1_score(y_test, y_pred, average='weighted')  # Check if current model is better than previous ones if f1 &gt; best_f1_score: best_f1_score = f1 best_params = { 'n_estimators': n_estimators, 'max_depth': max_depth, 'min_samples_split': min_samples_split,</pre>
---	--

		<pre> 'min_samples_leaf': min_samples_leaf }  # Train final RandomForestClassifier with best parameters best_rf_clf = RandomForestClassifier(n_estimators=best_params['n_estimators'], max_depth=best_params['max_depth'], min_samples_split=best_params['min_samples_split'], min_samples_leaf=best_params['min_samples_leaf'], n_jobs=-1, class_weight='balanced', random_state=42)  best_rf_clf.fit(X_train_scaled, y_train_res)  # Predict on test data with the best model y_pred = best_rf_clf.predict(X_test_scaled)  # Confusion Matrix conf_matrix = confusion_matrix(y_test, y_pred) tn, fp, fn, tp = conf_matrix.ravel()  # Calculate Metrics accuracy = accuracy_score(y_test, y_pred) f1 = f1_score(y_test, y_pred, average='weighted') precision = precision_score(y_test, y_pred, average='weighted') recall = recall_score(y_test, y_pred, average='weighted') specificity = tn / (tn + fp)  # Calculate Loss Cost cost_fp = 1000 cost_fn = 5000 loss_cost = cost_fp * fp + cost_fn * fn  print("""Best Parameters: """, best_params) print("""Best F1 Weighted Score: """, best_f1_score)  print("""Accuracy Score: """, accuracy) print("""F1 Score: """, f1) print("""Precision: """, precision) print("""Recall (Sensitivity): """, recall) print("""Specificity: """, specificity) print("""Loss Cost: """, loss_cost)  print("""Confusion Matrix: """) print(conf_matrix) print(f"""True Positives (TP): {tp}""") print(f"""True Negatives (TN): {tn}""") print(f"""False Positives (FP): {fp}""") print(f"""False Negatives (FN): {fn}""") </pre>
<b>After tuning- Model 2</b>	<pre> import pandas as pd import numpy as np from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing </pre>	<pre> # Define fixed set of features fixed_features = ['feature17', 'feature41', 'feature60', 'feature66', 'feature342', 'feature427', 'feature442', 'feature563', 'elapsed_time']  # Select only the fixed features X_train_selected = X_train[fixed_features] X_test_selected = X_test[fixed_features] </pre>

<pre>import MinMaxScaler from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score</pre>	<pre># SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected)  # Define parameter grid for RandomForestClassifier param_grid = { 'n_estimators': [50, 100, 200, 300], 'max_depth': [3, 5, 7, 10, 15], 'min_samples_split': [2, 5, 10, 15], 'min_samples_leaf': [1, 2, 4, 6] }  # Perform hyperparameter tuning without cross-validation best_f1_score = -1 best_params = None  for n_estimators in param_grid['n_estimators']: for max_depth in param_grid['max_depth']: for min_samples_split in param_grid['min_samples_split']: for min_samples_leaf in param_grid['min_samples_leaf']: # Train RandomForestClassifier with current parameters rf_clf = RandomForestClassifier(n_estimators=n_estimators, max_depth=max_depth, min_samples_split=min_samples_split, min_samples_leaf=min_samples_leaf, n_jobs=-1, class_weight='balanced', random_state=42)  rf_clf.fit(X_train_scaled, y_train_res)  # Predict on test data y_pred = rf_clf.predict(X_test_scaled)  # Calculate F1 score f1 = f1_score(y_test, y_pred, average='weighted')  # Check if current model is better than previous ones if f1 &gt; best_f1_score: best_f1_score = f1 best_params = { 'n_estimators': n_estimators, 'max_depth': max_depth, 'min_samples_split': min_samples_split, 'min_samples_leaf': min_samples_leaf }  # Train final RandomForestClassifier with best parameters best_rf_clf = RandomForestClassifier(n_estimators=best_params['n_estimators'], max_depth=best_params['max_depth'], min_samples_split=best_params['min_samples_split'], min_samples_leaf=best_params['min_samples_leaf'],</pre>
--	---

		<pre> n_jobs=-1, class_weight='balanced', random_state=42)  best_rf_clf.fit(X_train_scaled, y_train_res)  # Predict on test data with the best model y_pred = best_rf_clf.predict(X_test_scaled)  # Confusion Matrix conf_matrix = confusion_matrix(y_test, y_pred) tn, fp, fn, tp = conf_matrix.ravel()  # Calculate Metrics accuracy = accuracy_score(y_test, y_pred) f1 = f1_score(y_test, y_pred, average='weighted') precision = precision_score(y_test, y_pred, average='weighted') recall = recall_score(y_test, y_pred, average='weighted') specificity = tn / (tn + fp)  # Calculate Loss Cost cost_fp = 1000 cost_fn = 5000 loss_cost = cost_fp * fp + cost_fn * fn  print("Best Parameters:", best_params) print("Best F1 Weighted Score:", best_f1_score)  print("Accuracy Score: ", accuracy) print("F1 Score: ", f1) print("Precision: ", precision) print("Recall (Sensitivity): ", recall) print("Specificity: ", specificity) print("Loss Cost: ", loss_cost)  print("Confusion Matrix:") print(conf_matrix) print(f"True Positives (TP): {tp}") print(f"True Negatives (TN): {tn}") print(f"False Positives (FP): {fp}") print(f"False Negatives (FN): {fn}") </pre>
<b>After tuning- Model 3</b>	<pre> import pandas as pd import numpy as np from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score from sklearn.naive_bayes </pre>	<pre> # Selecting specific features identified by Boruta selected_features = ['feature60', 'feature65', 'feature66', 'feature342', 'feature351', 'feature478', 'feature540', 'feature563']  # Update X_train and X_test with selected features X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features]  # SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected) </pre>

import GaussianNB

```
# Define parameter grid for GaussianNB
param_grid = {
'var_smoothing': np.logspace(0,-9, num=100)
}

# Perform parameter tuning
best_f1_score = -1
best_params = None

for var_smoothing in param_grid['var_smoothing']:
# Train GaussianNB with current parameters
nb_clf = GaussianNB(var_smoothing=var_smoothing)
nb_clf.fit(X_train_scaled, y_train_res)

# Predict on test data
y_pred = nb_clf.predict(X_test_scaled)

# Calculate F1 score
f1 = f1_score(y_test, y_pred, average='weighted')

# Check if current model is better than previous ones
if f1 > best_f1_score:
best_f1_score = f1
best_params = {
'var_smoothing': var_smoothing
}

# Train final GaussianNB with best parameters
best_nb_clf =
GaussianNB(var_smoothing=best_params['var_smoothing'])
best_nb_clf.fit(X_train_scaled, y_train_res)

# Predict on test data with the best model
y_pred = best_nb_clf.predict(X_test_scaled)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = conf_matrix.ravel()

# Calculate Metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
specificity = tn / (tn + fp)

# Calculate Loss Cost
cost_fp = 1000
cost_fn = 5000
loss_cost = cost_fp * fp + cost_fn * fn

print("Best Parameters:", best_params)
print("Best F1 Weighted Score:", best_f1_score)

print("Accuracy Score: ", accuracy)
print("F1 Score: ", f1)
print("Precision: ", precision)
```



		<pre> print("Recall (Sensitivity): ", recall) print("Specificity: ", specificity) print("Loss Cost: ", loss_cost)  print("Confusion Matrix:") print(conf_matrix) print(f"True Positives (TP): {tp}") print(f"True Negatives (TN): {tn}") print(f"False Positives (FP): {fp}") print(f"False Negatives (FN): {fn}") </pre>
<b>Graphs-precision, f1, recall, lost... 3 models</b>	<pre> import numpy as np import matplotlib.pyplot as plt </pre>	<pre> # Data for three models before and after tuning models = ['Model 1', 'Model 2', 'Model 3'] metrics = ['Accuracy', 'F1 Score', 'Precision', 'Recall (Sensitivity)', 'Loss Cost'] metrics_values_before = np.array([ [0.8086734694, 0.8494935529, 0.9134543088, 0.8086734694, 119000], [0.9107142857, 0.2553191489, 0.2857142857, 0.2307692308, 115000], [0.5816326531, 0.6856787933, 0.8740304301, 0.5816326531, 228000] ]) metrics_values_after = np.array([ [0.8826530612, 0.8926764456, 0.904544958, 0.8826530612, 114000], [0.9285714286, 0.9184331797, 0.9126984127, 0.9285714286, 108000], [0.5816326531, 0.6856787933, 0.8740304301, 0.5816326531, 228000] ])  confusion_matrix_before = np.array([ [302, 64, 11, 15], [274, 19, 19, 2], [218, 148, 16, 10] ])  confusion_matrix_after = np.array([ [337, 29, 17, 9], [358, 8, 20, 6], [218, 148, 16, 10] ])  # Define colors for models colors = { 'Model 1': {'Before Tuning': 'lightgreen', 'After Tuning': 'green'}, 'Model 2': {'Before Tuning': 'mediumseagreen', 'After Tuning': 'seagreen'}, # Using mediumseagreen for Model 2 before tuning 'Model 3': {'Before Tuning': 'lightcoral', 'After Tuning': 'red'} }  # Function to plot side-by-side comparison of metrics def plot_metrics_comparison(metrics, metrics_values_before, metrics_values_after): num_metrics = len(metrics) fig, axes = plt.subplots(2, 3, figsize=(18, 10))  for i in range(num_metrics): row = i // 3 col = i % 3 ax = axes[row, col]  index = np.arange(len(models)) bar_width = 0.35 </pre>

```

before_vals = metrics_values_before[:, i]
after_vals = metrics_values_after[:, i]

for j, model in enumerate(models):
    ax.bar(index[j] - bar_width/2, before_vals[j], bar_width,
    label='Before Tuning', color=colors[model]['Before Tuning'], alpha=0.6)
    ax.bar(index[j] + bar_width/2, after_vals[j], bar_width,
    label='After Tuning', color=colors[model]['After Tuning'], alpha=1.0)

if metrics[i] == 'Loss Cost': # Adjusting format for Loss Cost metric
    ax.text(index[j] - bar_width/2, before_vals[j] + 5000, f'{before_vals[j]:.0f}',
    ha='center', va='bottom')
    ax.text(index[j] + bar_width/2, after_vals[j] + 5000, f'{after_vals[j]:.0f}',
    ha='center', va='bottom')
else:
    ax.text(index[j] - bar_width/2, before_vals[j] + 0.0005,
    f'{before_vals[j]:.3f}', ha='center', va='bottom')
    ax.text(index[j] + bar_width/2, after_vals[j] + 0.0005, f'{after_vals[j]:.3f}',
    ha='center', va='bottom')

ax.set_xlabel('Models')
ax.set_ylabel(metrics[i])
ax.set_title(f'{metrics[i]} Comparison')
ax.set_xticks(index)
ax.set_xticklabels(models)

# Add legend outside the plot with adjusted parameters
handles = [
    plt.Rectangle((0,0),1,1, color='lightgreen', alpha=0.6),
    plt.Rectangle((0,0),1,1, color='green'),
    plt.Rectangle((0,0),1,1, color='mediumseagreen'), # Using
    mediumseagreen for Model 2 before tuning
    plt.Rectangle((0,0),1,1, color='seagreen'),
    plt.Rectangle((0,0),1,1, color='lightcoral', alpha=0.6),
    plt.Rectangle((0,0),1,1, color='red')
]
labels = [
    'Model 1 - Before Tuning', 'Model 1 - After Tuning',
    'Model 2 - Before Tuning', 'Model 2 - After Tuning',
    'Model 3 - Before Tuning', 'Model 3 - After Tuning'
]
fig.legend(handles=handles, labels=labels, loc='upper center',
bbox_to_anchor=(0.5, 1.15), ncol=3, fontsize='large') # Adjusted
bbox_to_anchor and fontsize

# Adjust subplot spacing
plt.subplots_adjust(bottom=0.15) # Increase bottom padding for the
legend

plt.tight_layout()
plt.show()

# Function to plot side-by-side comparison of confusion matrix
components
def plot_confusion_matrix_comparison(confusion_matrix_before,
confusion_matrix_after):

```

		<pre> metrics_names = ['TN', 'FP', 'FN', 'TP'] num_metrics = len(metrics_names) fig, axes = plt.subplots(1, num_metrics, figsize=(18, 6))  for i, metric_name in enumerate(metrics_names):     ax = axes[i]      index = np.arange(len(models))     bar_width = 0.35      before_vals = confusion_matrix_before[:, i]     after_vals = confusion_matrix_after[:, i]      for j, model in enumerate(models):         ax.bar(index[j] - bar_width/2, before_vals[j], bar_width,             label='Before Tuning', color=colors[model]['Before Tuning'], alpha=0.6)         ax.bar(index[j] + bar_width/2, after_vals[j], bar_width,             label='After Tuning', color=colors[model]['After Tuning'], alpha=1.0)      ax.text(index[j] - bar_width/2, before_vals[j] + 0.0005,         f'{before_vals[j]:.0f}', ha='center', va='bottom')     ax.text(index[j] + bar_width/2, after_vals[j] + 0.0005, f'{after_vals[j]:.0f}',         ha='center', va='bottom')      ax.set_xlabel('Models')     ax.set_ylabel(metric_name)     ax.set_title(f'{metric_name} Comparison')     ax.set_xticks(index)     ax.set_xticklabels(models)      # Add legend outside the plot with adjusted parameters     handles = [         plt.Rectangle((0,0),1,1, color='lightgreen', alpha=0.6),         plt.Rectangle((0,0),1,1, color='green'),         plt.Rectangle((0,0),1,1, color='mediumseagreen'), # Using         mediumseagreen for Model 2 before tuning         plt.Rectangle((0,0),1,1, color='lightcoral', alpha=0.6),         plt.Rectangle((0,0),1,1, color='red')     ]     labels = [         'Model 1 - Before Tuning', 'Model 1 - After Tuning',         'Model 2 - Before Tuning', 'Model 2 - After Tuning',         'Model 3 - Before Tuning', 'Model 3 - After Tuning'     ]      plt.tight_layout()     plt.show()      # Plotting metrics comparison     plot_metrics_comparison(metrics, metrics_values_before[:, :5],         metrics_values_after[:, :5])      # Plotting confusion matrix components comparison     plot_confusion_matrix_comparison(confusion_matrix_before,         confusion_matrix_after) </pre>
<b>Trade off 1- Rates - 3</b>	<pre> import numpy as np import matplotlib.pyplot as plt </pre>	<pre> # Define the values for Model 1, Model 2, and Model 3 after tuning models = ['Model 1', 'Model 2', 'Model 3'] </pre>

models	<pre> TP = np.array([9, 6, 10]) TN = np.array([337, 358, 218]) FP = np.array([29, 8, 148]) FN = np.array([17, 20, 16])  # Calculate rates total_positives = TP + FN total_negatives = TN + FP  TP_rate = TP / total_positives FP_rate = FP / total_negatives FN_rate = FN / total_positives TN_rate = TN / total_negatives  # Plotting the comparison graph rates = ['TP Rate', 'FP Rate', 'FN Rate', 'TN Rate']  # Transpose the rates for plotting rate_values = np.array([TP_rate, FP_rate, FN_rate, TN_rate])  plt.figure(figsize=(10, 6))  # Plot each line separately for each model for i, model in enumerate(models): plt.plot(rates, rate_values[:, i], marker='o', linestyle='-', label=model)  # Adding labels and title plt.xlabel('Rates') plt.ylabel('Value') plt.title('Comparison of TP, FP, FN, TN Rates After Tuning for Model 1, Model 2, and Model 3') plt.xticks(rotation=45) plt.legend()  plt.grid(True) plt.tight_layout() plt.show() </pre>
--------	--

## 6.6 Cross Validation and Evaluation

Modeling and Evaluation		
Purpose	Library	Query
<b>Kfold- Model 1</b>	import pandas as pd import numpy as np from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split, StratifiedKFold from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score	# Assuming 'merged_df' is your DataFrame with the relevant data X = merged_df.drop('feature591', axis=1) # Replace 'feature591' with your actual target column name y = merged_df['feature591']  # Initialize k-fold cross-validation kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)  # Initialize lists to store evaluation metrics across folds accuracy_scores = [] f1_scores = [] precision_scores = [] recall_scores = [] specificity_scores = [] loss_costs = []  # Iterate over each fold for fold_idx, (train_idx, test_idx) in enumerate(kfold.split(X, y)): print(f"\nFold {fold_idx + 1}:")  # Split data into training and testing sets for this fold X_train, X_test = X.iloc[train_idx], X.iloc[test_idx] y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]  # Boruta Feature Selection rf_clf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5, random_state=42) boruta_selector = BorutaPy(rf_clf, n_estimators='auto', verbose=0, random_state=42) # Set verbose to 0 boruta_selector.fit(X_train.values, y_train)  # Selected features selected_features = X_train.columns[boruta_selector.support_].tolist() print("Selected Features: ", selected_features)  # Update X_train and X_test with selected features X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features]  # SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected)  # Define parameter grid for RandomForestClassifier param_grid = { 'n_estimators': [50, 100, 200, 300], 'max_depth': [3, 5, 7, 10, 15],

```

'min_samples_split': [2, 5, 10, 15],
'min_samples_leaf': [1, 2, 4, 6]
}

# Perform hyperparameter tuning without cross-validation
best_f1_score = -1
best_params = None

for n_estimators in param_grid['n_estimators']:
for max_depth in param_grid['max_depth']:
for min_samples_split in param_grid['min_samples_split']:
for min_samples_leaf in param_grid['min_samples_leaf']:
# Train RandomForestClassifier with current parameters
rf_clf = RandomForestClassifier(n_estimators=n_estimators,
max_depth=max_depth,
min_samples_split=min_samples_split,
min_samples_leaf=min_samples_leaf,
n_jobs=-1,
class_weight='balanced',
random_state=42)

rf_clf.fit(X_train_scaled, y_train_res)

# Predict on test data
y_pred = rf_clf.predict(X_test_scaled)

# Calculate F1 score
f1 = f1_score(y_test, y_pred, average='weighted')

# Check if current model is better than previous ones
if f1 > best_f1_score:
best_f1_score = f1
best_params = {
'n_estimators': n_estimators,
'max_depth': max_depth,
'min_samples_split': min_samples_split,
'min_samples_leaf': min_samples_leaf
}

# Train final RandomForestClassifier with best parameters
best_rf_clf =
RandomForestClassifier(n_estimators=best_params['n_estimators'],
max_depth=best_params['max_depth'],
min_samples_split=best_params['min_samples_split'],
min_samples_leaf=best_params['min_samples_leaf'],
n_jobs=-1,
class_weight='balanced',
random_state=42)

best_rf_clf.fit(X_train_scaled, y_train_res)

# Predict on test data with the best model
y_pred = best_rf_clf.predict(X_test_scaled)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = conf_matrix.ravel()

```

		<pre> # Calculate Metrics accuracy = accuracy_score(y_test, y_pred) f1 = f1_score(y_test, y_pred, average='weighted') precision = precision_score(y_test, y_pred, average='weighted') recall = recall_score(y_test, y_pred, average='weighted') specificity = tn / (tn + fp)  # Calculate Loss Cost cost_fp = 1000 cost_fn = 5000 loss_cost = cost_fp * fp + cost_fn * fn  # Append scores to lists accuracy_scores.append(accuracy) f1_scores.append(f1) precision_scores.append(precision) recall_scores.append(recall) specificity_scores.append(specificity) loss_costs.append(loss_cost)  # Print evaluation metrics for the fold print("\nEvaluation Metrics for Fold:") print("Best Parameters:", best_params) print("Best F1 Weighted Score:", best_f1_score) print("Accuracy Score: ", accuracy) print("F1 Score: ", f1) print("Precision: ", precision) print("Recall (Sensitivity): ", recall) print("Specificity: ", specificity) print("Loss Cost: ", loss_cost)  print("Confusion Matrix:") print(conf_matrix) print(f"True Positives (TP): {tp}") print(f"True Negatives (TN): {tn}") print(f"False Positives (FP): {fp}") print(f"False Negatives (FN): {fn}")  # Print average scores across all folds print("\nAverage Metrics Across All Folds:") print("Average Accuracy Score: ", np.mean(accuracy_scores)) print("Average F1 Score: ", np.mean(f1_scores)) print("Average Precision: ", np.mean(precision_scores)) print("Average Recall (Sensitivity): ", np.mean(recall_scores)) print("Average Specificity: ", np.mean(specificity_scores)) print("Average Loss Cost: ", np.mean(loss_costs)) </pre>
<b>Kfold- Model 2</b>	<pre> import pandas as pd import numpy as np from sklearn.impute import KNNImputer from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler </pre>	<pre> # Initialize k-fold cross-validation kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)  # Initialize lists to store evaluation metrics across folds accuracy_scores = [] f1_scores = [] precision_scores = [] recall_scores = [] specificity_scores = [] loss_costs = []  # Assuming X_train, y_train, X_test, y_test are already defined </pre>

<pre> from sklearn.model_selection import StratifiedKFold from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score </pre>	<pre> X_train = X_train_cleaned[['feature17', 'feature41', 'feature60', 'feature66', 'feature342', 'feature427', 'feature442', 'feature563', 'elapsed_time']].copy() y_train = y_train.copy() X_test = X_test_cleaned[['feature17', 'feature41', 'feature60', 'feature66', 'feature342', 'feature427', 'feature442', 'feature563', 'elapsed_time']].copy() y_test = y_test.copy()  # Perform KNN imputation within each fold for train_index, test_index in kfold.split(X_train, y_train): X_train_fold, X_val_fold = X_train.iloc[train_index], X_train.iloc[test_index] y_train_fold, y_val_fold = y_train.iloc[train_index], y_train.iloc[test_index]  # KNN imputation on training fold imputer = KNNImputer(n_neighbors=5) X_train_imputed = imputer.fit_transform(X_train_fold) X_val_imputed = imputer.transform(X_val_fold)  # SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_imputed, y_train_fold)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_val_scaled = scaler.transform(X_val_imputed)  # Define parameter grid for RandomForestClassifier param_grid = { 'n_estimators': [50, 100, 200, 300], 'max_depth': [3, 5, 7, 10, 15], 'min_samples_split': [2, 5, 10, 15], 'min_samples_leaf': [1, 2, 4, 6] }  # Perform hyperparameter tuning without cross-validation best_f1_score = -1 best_params = None  for n_estimators in param_grid['n_estimators']: for max_depth in param_grid['max_depth']: for min_samples_split in param_grid['min_samples_split']: for min_samples_leaf in param_grid['min_samples_leaf']: # Train RandomForestClassifier with current parameters rf_clf = RandomForestClassifier(n_estimators=n_estimators, max_depth=max_depth, min_samples_split=min_samples_split, min_samples_leaf=min_samples_leaf, n_jobs=-1, class_weight='balanced', random_state=42) rf_clf.fit(X_train_scaled, y_train_res)  # Predict on validation data </pre>
--	--



```

y_pred = rf_clf.predict(X_val_scaled)

# Calculate F1 score
f1 = f1_score(y_val_fold, y_pred, average='weighted')

# Check if current model is better than previous ones
if f1 > best_f1_score:
    best_f1_score = f1
    best_params = {
        'n_estimators': n_estimators,
        'max_depth': max_depth,
        'min_samples_split': min_samples_split,
        'min_samples_leaf': min_samples_leaf
    }

# Train final RandomForestClassifier with best parameters on full
training fold
best_rf_clf =
RandomForestClassifier(n_estimators=best_params['n_estimators'],
max_depth=best_params['max_depth'],
min_samples_split=best_params['min_samples_split'],
min_samples_leaf=best_params['min_samples_leaf'],
n_jobs=-1,
class_weight='balanced',
random_state=42)

best_rf_clf.fit(X_train_scaled, y_train_res)

# Predict on test data with the best model
X_test_imputed = imputer.transform(X_test) # Impute test data
X_test_scaled = scaler.transform(X_test_imputed) # Scale test data
y_pred = best_rf_clf.predict(X_test_scaled)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = conf_matrix.ravel()

# Calculate Metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
specificity = tn / (tn + fp)

# Calculate Loss Cost
cost_fp = 1000
cost_fn = 5000
loss_cost = cost_fp * fp + cost_fn * fn

# Append scores to lists
accuracy_scores.append(accuracy)
f1_scores.append(f1)
precision_scores.append(precision)
recall_scores.append(recall)
specificity_scores.append(specificity)
loss_costs.append(loss_cost)

# Print evaluation metrics for the fold

```

		<pre> print("\nEvaluation Metrics for Fold:") print("Best Parameters:", best_params) print("Best F1 Weighted Score:", best_f1_score) print("Accuracy Score: ", accuracy) print("F1 Score: ", f1) print("Precision: ", precision) print("Recall (Sensitivity): ", recall) print("Specificity: ", specificity) print("Loss Cost: ", loss_cost)  print("Confusion Matrix:") print(conf_matrix) print(f"True Positives (TP): {tp}") print(f"True Negatives (TN): {tn}") print(f"False Positives (FP): {fp}") print(f"False Negatives (FN): {fn}")  # Print average scores across all folds print("\nAverage Metrics Across All Folds:") print("Average Accuracy Score: ", np.mean(accuracy_scores)) print("Average F1 Score: ", np.mean(f1_scores)) print("Average Precision: ", np.mean(precision_scores)) print("Average Recall (Sensitivity): ", np.mean(recall_scores)) print("Average Specificity: ", np.mean(specificity_scores)) print("Average Loss Cost: ", np.mean(loss_costs)) </pre>
<b>Kfold Volatility comparison for Model 1 and 2</b>	import matplotlib.pyplot as plt	<pre> # Loss costs for each fold fold_loss_costs = [114000, 99000, 85000, 103000, 94000]  # Calculate average loss cost avg_loss_cost = 99000  # Labels for folds fold_labels = ['Fold 1', 'Fold 2', 'Fold 3', 'Fold 4', 'Fold 5']  # Plotting the loss costs plt.figure(figsize=(10, 6)) plt.bar(fold_labels, fold_loss_costs, color='skyblue', label='Loss Cost per Fold') plt.axhline(y=avg_loss_cost, color='orange', linestyle='--', label='Average Loss Cost') plt.xlabel('Folds') plt.ylabel('Loss Cost') plt.title('Loss Cost Comparison Across Folds') plt.ylim(0, max(fold_loss_costs) * 1.2) # Adjust ylim for better visualization plt.legend() plt.grid(True) plt.tight_layout()  # Adding values above bars for i, v in enumerate(fold_loss_costs): plt.text(i, v + 5000, str(v), ha='center', va='bottom', fontsize=10)  # Adding average value annotation plt.text(len(fold_loss_costs) - 0.5, avg_loss_cost + 5000, f'Avg: {avg_loss_cost}', ha='center', va='bottom', fontsize=10, color='orange')  plt.show() </pre>

Trade off 2- Learning Curve for Model 1		
Before tuning	<pre>import pandas as pd import numpy as np import matplotlib.pyplot as plt from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split, learning_curve, validation_curve from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score</pre>	<pre># Assuming 'merged_df' is your DataFrame with the relevant data # Replace 'merged_df' with your actual DataFrame containing the data # merged_df = pd.read_csv('your_data.csv') # Replace with your data loading code X = merged_df.drop('feature591', axis=1) # Replace 'feature591' with your actual target column name y = merged_df['feature591']  # Split data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)  # Boruta Feature Selection rf_clf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5, random_state=42) boruta_selector = BorutaPy(rf_clf, n_estimators='auto', verbose=0, random_state=42) # Set verbose to 0 boruta_selector.fit(X_train.values, y_train)  # Selected features selected_features = X_train.columns[boruta_selector.support_].tolist() print("Selected Features: ", selected_features)  # Update X_train and X_test with selected features X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features]  # SMOTE Balancing on training data smote = SMOTE(random_state=42) X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)  # Min-Max Scaling scaler = MinMaxScaler() X_train_scaled = scaler.fit_transform(X_train_res) X_test_scaled = scaler.transform(X_test_selected)  # Define parameter grid for RandomForestClassifier param_grid = { 'n_estimators': [50, 100, 200, 300], 'max_depth': [3, 5, 7, 10, 15], 'min_samples_split': [2, 5, 10, 15], 'min_samples_leaf': [1, 2, 4, 6] }  # Perform hyperparameter tuning without cross-validation best_f1_score = -1 best_params = None  for n_estimators in param_grid['n_estimators']: for max_depth in param_grid['max_depth']: for min_samples_split in param_grid['min_samples_split']: for min_samples_leaf in param_grid['min_samples_leaf']: # Train RandomForestClassifier with current parameters rf_clf = RandomForestClassifier(n_estimators=n_estimators,</pre>

```

max_depth=max_depth,
min_samples_split=min_samples_split,
min_samples_leaf=min_samples_leaf,
n_jobs=-1,
class_weight='balanced',
random_state=42)

rf_clf.fit(X_train_scaled, y_train_res)

# Predict on test data
y_pred = rf_clf.predict(X_test_scaled)

# Calculate F1 score
f1 = f1_score(y_test, y_pred, average='weighted')

# Check if current model is better than previous ones
if f1 > best_f1_score:
    best_f1_score = f1
    best_params = {
        'n_estimators': n_estimators,
        'max_depth': max_depth,
        'min_samples_split': min_samples_split,
        'min_samples_leaf': min_samples_leaf
    }

# Train final RandomForestClassifier with best parameters
best_rf_clf =
RandomForestClassifier(n_estimators=best_params['n_estimators'],
max_depth=best_params['max_depth'],
min_samples_split=best_params['min_samples_split'],
min_samples_leaf=best_params['min_samples_leaf'],
n_jobs=-1,
class_weight='balanced',
random_state=42)

best_rf_clf.fit(X_train_scaled, y_train_res)

# Predict on test data with the best model
y_pred = best_rf_clf.predict(X_test_scaled)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = conf_matrix.ravel()

# Calculate Metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
specificity = tn / (tn + fp)

# Calculate Loss Cost
cost_fp = 1000
cost_fn = 5000
loss_cost = cost_fp * fp + cost_fn * fn

print("Best Parameters:", best_params)
print("Best F1 Weighted Score:", best_f1_score)

```

```

print("Accuracy Score: ", accuracy)
print("F1 Score: ", f1)
print("Precision: ", precision)
print("Recall (Sensitivity): ", recall)
print("Specificity: ", specificity)
print("Loss Cost: ", loss_cost)

print("Confusion Matrix:")
print(conf_matrix)
print(f"True Positives (TP): {tp}")
print(f"True Negatives (TN): {tn}")
print(f"False Positives (FP): {fp}")
print(f"False Negatives (FN): {fn}")

# Plotting the learning curve
def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
scoring='f1_weighted', n_jobs=-1, train_sizes=np.linspace(.1, 1.0, 5)):
plt.figure()
plt.title(title)
if ylim is not None:
plt.ylim(*ylim)
plt.xlabel("Training examples")
plt.ylabel("Score")
train_sizes, train_scores, test_scores = learning_curve(
estimator, X, y, cv=cv, scoring=scoring, n_jobs=n_jobs,
train_sizes=train_sizes)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
plt.grid()

plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
train_scores_mean + train_scores_std, alpha=0.1,
color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
label="Cross-validation score")

plt.legend(loc="best")
return plt

title = "Learning Curves (RandomForestClassifier)"
plot_learning_curve(best_rf_clf, title, X_train_scaled, y_train_res, cv=5,
scoring='f1_weighted')
plt.show()

# Plotting the validation curve
def plot_validation_curve(estimator, title, X, y, param_name,
param_range, ylim=None, cv=None,
scoring="f1_weighted", n_jobs=-1):
plt.figure()
plt.title(title)
if ylim is not None:

```

		<pre> plt.ylim(*ylim) plt.xlabel(param_name) plt.ylabel("Score") train_scores, test_scores = validation_curve( estimator, X, y, param_name=param_name, param_range=param_range, cv=cv, scoring=scoring, n_jobs=n_jobs) train_scores_mean = np.mean(train_scores, axis=1) train_scores_std = np.std(train_scores, axis=1) test_scores_mean = np.mean(test_scores, axis=1) test_scores_std = np.std(test_scores, axis=1) plt.grid()  plt.semilogx(param_range, train_scores_mean, label="Training score", color="darkorange", lw=2) plt.fill_between(param_range, train_scores_mean - train_scores_std, train_scores_mean + train_scores_std, alpha=0.2, color="darkorange", lw=2) plt.semilogx(param_range, test_scores_mean, label="Cross-validation score", color="navy", lw=2) plt.fill_between(param_range, test_scores_mean - test_scores_std, test_scores_mean + test_scores_std, alpha=0.2, color="navy", lw=2) plt.legend(loc="best") return plt  title = "Validation Curve (RandomForestClassifier)" param_name = "n_estimators" param_range = [50, 100, 200, 300] # Specify the range of n_estimators plot_validation_curve(best_rf_clf, title, X_train_scaled, y_train_res, param_name=param_name, param_range=param_range, cv=5, scoring="f1_weighted") plt.show() </pre>
<b>After tuning</b>	<pre> import pandas as pd import numpy as np import matplotlib.pyplot as plt from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split, learning_curve from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score </pre>	<pre> # Assuming 'merged_df' is your DataFrame with the relevant data X = merged_df.drop('feature591', axis=1) # Replace 'feature591' with your actual target column name y = merged_df['feature591']  # Split data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)  # Boruta Feature Selection rf_clf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5, random_state=42) boruta_selector = BorutaPy(rf_clf, n_estimators='auto', verbose=0, random_state=42) # Set verbose to 0 boruta_selector.fit(X_train.values, y_train)  # Selected features selected_features = X_train.columns[boruta_selector.support_].tolist() print("Selected Features: ", selected_features)  # Update X_train and X_test with selected features X_train_selected = X_train[selected_features] X_test_selected = X_test[selected_features] </pre>

```

# SMOTE Balancing on training data
smote = SMOTE(random_state=42)
X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)

# Min-Max Scaling
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train_res)
X_test_scaled = scaler.transform(X_test_selected)

# Train Random Forest Classifier
rf_clf.fit(X_train_scaled, y_train_res)
y_pred = rf_clf.predict(X_test_scaled)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = conf_matrix.ravel()

# Calculate Metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
specificity = tn / (tn + fp)

# Calculate Loss Cost
cost_fp = 1000
cost_fn = 5000
loss_cost = cost_fp * fp + cost_fn * fn

print("Accuracy Score: ", accuracy)
print("F1 Score: ", f1)
print("Precision: ", precision)
print("Recall (Sensitivity): ", recall)
print("Specificity: ", specificity)
print("Loss Cost: ", loss_cost)

print("Confusion Matrix:")
print(conf_matrix)
print(f"True Positives (TP): {tp}")
print(f"True Negatives (TN): {tn}")
print(f"False Positives (FP): {fp}")
print(f"False Negatives (FN): {fn}")

# Learning Curve
train_sizes, train_scores, test_scores = learning_curve(
    rf_clf, X_train_scaled, y_train_res, cv=5, scoring='f1_weighted',
    n_jobs=-1,
    train_sizes=np.linspace(0.1, 1.0, 10))

# Calculate mean and standard deviation of training scores and test scores
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

# Plotting the learning curve
plt.figure(figsize=(10, 6))

```

		<pre> plt.title("Learning Curve (RandomForestClassifier)") plt.xlabel("Training examples") plt.ylabel("Score")  plt.grid() plt.fill_between(train_sizes, train_scores_mean - train_scores_std, train_scores_mean + train_scores_std, alpha=0.1, color="r") plt.fill_between(train_sizes, test_scores_mean - test_scores_std, test_scores_mean + test_scores_std, alpha=0.1, color="g") plt.plot(train_sizes, train_scores_mean, 'o-', color="r", label="Training score") plt.plot(train_sizes, test_scores_mean, 'o-', color="g", label="Cross-validation score") plt.legend(loc="best")  plt.show() </pre>
<b>After CV</b>	<pre> import pandas as pd import numpy as np from boruta import BorutaPy from sklearn.ensemble import RandomForestClassifier from imblearn.over_sampling import SMOTE from sklearn.preprocessing import MinMaxScaler from sklearn.model_selection import train_test_split, StratifiedKFold, learning_curve from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score import matplotlib.pyplot as plt </pre>	<pre> # Assuming 'merged_df' is your DataFrame with the relevant data X = merged_df.drop('feature591', axis=1) # Replace 'feature591' with your actual target column name y = merged_df['feature591']  # Initialize k-fold cross-validation kfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)  # Initialize lists to store evaluation metrics across folds accuracy_scores = [] f1_scores = [] precision_scores = [] recall_scores = [] specificity_scores = [] loss_costs = []  # Initialize lists to store learning curve data train_sizes_all = [] train_scores_mean_all = [] train_scores_std_all = [] test_scores_mean_all = [] test_scores_std_all = []  # Iterate over each fold for fold_idx, (train_idx, test_idx) in enumerate(kfold.split(X, y)):     print(f"\nFold {fold_idx + 1}:")      # Split data into training and testing sets for this fold     X_train, X_test = X.iloc[train_idx], X.iloc[test_idx]     y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]      # Boruta Feature Selection     rf_clf = RandomForestClassifier(n_jobs=-1, class_weight='balanced', max_depth=5, random_state=42)     boruta_selector = BorutaPy(rf_clf, n_estimators='auto', verbose=0, random_state=42) # Set verbose to 0     boruta_selector.fit(X_train.values, y_train)      # Selected features     selected_features = X_train.columns[boruta_selector.support_].tolist()     print("Selected Features: ", selected_features) </pre>



```

# Update X_train and X_test with selected features
X_train_selected = X_train[selected_features]
X_test_selected = X_test[selected_features]

# SMOTE Balancing on training data
smote = SMOTE(random_state=42)
X_train_res, y_train_res = smote.fit_resample(X_train_selected, y_train)

# Min-Max Scaling
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train_res)
X_test_scaled = scaler.transform(X_test_selected)

# Define parameter grid for RandomForestClassifier
param_grid = {
    'n_estimators': [50, 100, 200, 300],
    'max_depth': [3, 5, 7, 10, 15],
    'min_samples_split': [2, 5, 10, 15],
    'min_samples_leaf': [1, 2, 4, 6]
}

# Perform hyperparameter tuning without cross-validation
best_f1_score = -1
best_params = None

for n_estimators in param_grid['n_estimators']:
    for max_depth in param_grid['max_depth']:
        for min_samples_split in param_grid['min_samples_split']:
            for min_samples_leaf in param_grid['min_samples_leaf']:
                # Train RandomForestClassifier with current parameters
                rf_clf = RandomForestClassifier(n_estimators=n_estimators,
                                                max_depth=max_depth,
                                                min_samples_split=min_samples_split,
                                                min_samples_leaf=min_samples_leaf,
                                                n_jobs=-1,
                                                class_weight='balanced',
                                                random_state=42)

                rf_clf.fit(X_train_scaled, y_train_res)

                # Predict on test data
                y_pred = rf_clf.predict(X_test_scaled)

                # Calculate F1 score
                f1 = f1_score(y_test, y_pred, average='weighted')

                # Check if current model is better than previous ones
                if f1 > best_f1_score:
                    best_f1_score = f1
                    best_params = {
                        'n_estimators': n_estimators,
                        'max_depth': max_depth,
                        'min_samples_split': min_samples_split,
                        'min_samples_leaf': min_samples_leaf
                    }

# Train final RandomForestClassifier with best parameters

```

```

best_rf_clf =
RandomForestClassifier(n_estimators=best_params['n_estimators'],
max_depth=best_params['max_depth'],
min_samples_split=best_params['min_samples_split'],
min_samples_leaf=best_params['min_samples_leaf'],
n_jobs=-1,
class_weight='balanced',
random_state=42)

best_rf_clf.fit(X_train_scaled, y_train_res)

# Predict on test data with the best model
y_pred = best_rf_clf.predict(X_test_scaled)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = conf_matrix.ravel()

# Calculate Metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
specificity = tn / (tn + fp)

# Calculate Loss Cost
cost_fp = 1000
cost_fn = 5000
loss_cost = cost_fp * fp + cost_fn * fn

# Append scores to lists
accuracy_scores.append(accuracy)
f1_scores.append(f1)
precision_scores.append(precision)
recall_scores.append(recall)
specificity_scores.append(specificity)
loss_costs.append(loss_cost)

# Print evaluation metrics for the fold
print("\nEvaluation Metrics for Fold:")
print("Best Parameters:", best_params)
print("Best F1 Weighted Score:", best_f1_score)
print("Accuracy Score: ", accuracy)
print("F1 Score: ", f1)
print("Precision: ", precision)
print("Recall (Sensitivity): ", recall)
print("Specificity: ", specificity)
print("Loss Cost: ", loss_cost)

print("Confusion Matrix:")
print(conf_matrix)
print(f"True Positives (TP): {tp}")
print(f"True Negatives (TN): {tn}")
print(f"False Positives (FP): {fp}")
print(f"False Negatives (FN): {fn}")

# Calculate learning curve
train_sizes, train_scores, test_scores, fit_times, _ = learning_curve(

```

```

best_rf_clf, X_train_scaled, y_train_res, cv=kfold, n_jobs=-1,
scoring='f1_weighted', train_sizes=np.linspace(.1, 1.0, 5),
return_times=True)

# Store learning curve data
train_sizes_all.append(train_sizes)
train_scores_mean_all.append(np.mean(train_scores, axis=1))
train_scores_std_all.append(np.std(train_scores, axis=1))
test_scores_mean_all.append(np.mean(test_scores, axis=1))
test_scores_std_all.append(np.std(test_scores, axis=1))

# Print average scores across all folds
print("\nAverage Metrics Across All Folds:")
print("Average Accuracy Score: ", np.mean(accuracy_scores))
print("Average F1 Score: ", np.mean(f1_scores))
print("Average Precision: ", np.mean(precision_scores))
print("Average Recall (Sensitivity): ", np.mean(recall_scores))
print("Average Specificity: ", np.mean(specificity_scores))
print("Average Loss Cost: ", np.mean(loss_costs))

# Plot learning curve
plt.figure()
plt.title("Learning Curve")
plt.xlabel("Training examples")
plt.ylabel("Score")

plt.grid()

train_scores_mean_all = np.array(train_scores_mean_all)
train_scores_std_all = np.array(train_scores_std_all)
test_scores_mean_all = np.array(test_scores_mean_all)
test_scores_std_all = np.array(test_scores_std_all)

plt.fill_between(train_sizes_all[0], train_scores_mean_all.mean(axis=0) -
train_scores_std_all.mean(axis=0),
train_scores_mean_all.mean(axis=0) +
train_scores_std_all.mean(axis=0), alpha=0.1,
color="r")
plt.fill_between(train_sizes_all[0], test_scores_mean_all.mean(axis=0) -
test_scores_std_all.mean(axis=0),
test_scores_mean_all.mean(axis=0) +
test_scores_std_all.mean(axis=0), alpha=0.1, color="g")
plt.plot(train_sizes_all[0], train_scores_mean_all.mean(axis=0), 'o-',
color="r",
label="Training score")
plt.plot(train_sizes_all[0], test_scores_mean_all.mean(axis=0), 'o-',
color="g",
label="Cross-validation score")

plt.legend(loc="best")
plt.show()

```